

---

# ADVANCED INFORMATION SYSTEM SECURITY

---

Hoping to get a better grade this time around.

**Fabio Lorenzato**

October 20, 2024

# Contents

<b>1</b>	<b>Transport Layer Security</b>	<b>3</b>
1.1	TLS session and connection . . . . .	4
1.2	TLS handshake protocol . . . . .	4
1.3	Achieving Data protection . . . . .	4
1.4	Relationship among keys and sessions . . . . .	5
1.5	Perfect Forward Secrecy . . . . .	6
1.6	The protocol . . . . .	6
1.6.1	Client Hello and Server Hello . . . . .	6
1.6.2	Cipher suite . . . . .	7
1.6.3	Certificates . . . . .	7
1.6.4	Key exchange . . . . .	8
1.6.5	Certificate verify . . . . .	8
1.6.6	Change cipher spec . . . . .	8
1.6.7	Finished message . . . . .	8
1.7	Setup Time . . . . .	9
1.8	TLS versions . . . . .	9
1.8.1	TLS 1.0 . . . . .	9
1.8.2	TLS 1.1 . . . . .	10
1.8.3	TLS 1.2 . . . . .	10
1.9	TLS attacks . . . . .	10
1.9.1	Heartbleed . . . . .	10
1.9.2	Bleichenbacher attack . . . . .	10
1.9.3	Other attacks against SSL/TLS . . . . .	11
1.10	ALPN extension . . . . .	12
1.11	TLS False Start . . . . .	12
1.12	The TLS downgrade problem . . . . .	13
1.12.1	TLS Fallback Signalling Cipher Suite Value (SCSV) . . . . .	13
1.13	TLS session tickets . . . . .	13
1.14	The Virtual Server Problem . . . . .	14
1.15	TLS 1.3 . . . . .	14
1.15.1	Key exchange . . . . .	14
1.15.2	Message protection . . . . .	15
1.15.3	Digital signature . . . . .	15
1.15.4	Ciphersuites . . . . .	15
1.15.5	EdDSA . . . . .	15

1.15.6	Other improvements . . . . .	16
1.15.7	HKDF in TLS 1.3 . . . . .	16
1.15.8	TLS-1.3 handshake . . . . .	17
1.16	TLS and PKI . . . . .	19
1.17	TLS and certificate status . . . . .	20
1.17.1	Pushed CRL . . . . .	20
1.18	OCSP Stapling . . . . .	20
1.18.1	OCSP Must Staple . . . . .	22
1.19	Questions and answers . . . . .	23
<b>2</b>	<b>SSH</b>	<b>25</b>
2.1	Architecture . . . . .	25
2.1.1	SSH Transport Layer Protocol . . . . .	25
2.1.2	Encryption . . . . .	30
2.1.3	MAC . . . . .	30
2.1.4	Peer authentication . . . . .	31
2.2	Port forwarding . . . . .	31
2.2.1	Local port forwarding . . . . .	31
2.2.2	Remote port forwarding . . . . .	32
2.2.3	Causes of insecurity . . . . .	32
2.3	Attacks on SSH . . . . .	33
2.3.1	BothanSpy . . . . .	33
2.3.2	Gyr Falcon . . . . .	33
2.3.3	Brute force attack . . . . .	34
2.4	Main Applications of SSH . . . . .	35
<b>3</b>	<b>The X.509 standard and PKI</b>	<b>36</b>
3.1	Public-key certificate (PKC) . . . . .	36
3.1.1	Key Generation in Public Key Cryptography (PKC) . . . . .	36
3.1.2	Certification architecture . . . . .	37
3.1.3	Certificate generation . . . . .	38
3.2	X.509 certificates . . . . .	39
3.2.1	PKC Scope . . . . .	39
3.2.2	Certificate Policy (CP) and Certification Practice Statement (CPS) . . . . .	40
3.2.3	X.500 Directory Service . . . . .	40
3.2.4	RFC-1422 . . . . .	41
3.3	X.509 Version 3 . . . . .	42
3.3.1	Base syntax . . . . .	43
3.3.2	Critical Extensions . . . . .	44
3.3.3	Public Extensions . . . . .	45
3.3.4	Key and policy information . . . . .	45
3.3.5	Certificate subject and issuer attributes . . . . .	48
3.3.6	Certificate path constraints . . . . .	50
3.3.7	CRL distribution point . . . . .	51
3.3.8	Private Extensions . . . . .	51

# Chapter 1

## Transport Layer Security

TLS, or Transport Layer Security, was originally proposed by Netscape in 1995 as a way to secure communications between a web browser and a web server. It is the successor to SSL, or Secure Sockets Layer, which was first introduced by Netscape in 1995. The two terms are often used interchangeably, but TLS is the more modern and secure protocol.

The main goal of SSL was to create secure network channel, almost at session level(4.5), between two parties, to provide some security services that neither TCP nor IP provides:

- **peer authentication** based on asymmetric challenge-response authentication(the challenge for the service is implicit, while for the client is explicit). Server authentication is always compulsory, while client authentication is optional and requested by the server.
- **message confidentiality** base on symmetric encryption
- **message integrity** and authentication based on MAC computed on the trasmitted data
- **replay, filtering and reordering attack protection** using implicit record numbers(the correct order of transmission is provided by TCP, for this reason the number is implicit). This number is used also in the MAC computation.

You can see the TLS packet structure in figure 1.1. The TLS handshake protocol is used to establish a new session or reestablish an existing session. The TLS change cipher spec protocol is used to trigger the change of the algorithms to be used for message protection, or most notably to pass from the previous unprotected session to a protected one. The TLS alert protocol is used to signal errors or signal the end of the connection. The TLS record protocol contains the generic protocols informations and its content depend of the state of the connection and the protocol it is tunneling.




Figure 1.1: TLS packet structure.

## 1.1 TLS session and connection

It is important to make a clear distinction between TLS session and connections.

**TLS sessions** a **logical association** between client and server, created via an handshake protocol and its shared between different TLS connections(1:N).

**TLS connections** are a **transient TLS channel** between client and server, which means that each connection is associated with only one specific TLS session(1:1).




Figure 1.2: TLS session and connection.

## 1.2 TLS handshake protocol

The TLS handshake protocol is used to establish a new session or reestablish an existing session. It's a critical part of the TLS protocol, because the channel pass from an unprotected state to a protected one. During this phase the two parts agree on a set of algorithms for confidentiality and integrity, exchange random numbers between the client and the server to be used for the subsequent generation of the keys, establish a symmetric key by means of public key operations (originally RSA and DHKE, but nowadays the elliptic curve versions of algorithms are used) and negotiate the session-id and exchange the necessary public keys certificates for the asymmetric challenge-response authentication.

## 1.3 Achieving Data protection

Data protection is achieved by using symmetric encryption algorithms to encrypt the data and Message Authentication Codes(MAC) to ensure the integrity of the data and the authentication of the sender.

Figure 1.3 shows how the data protection is achieved in TLS using authenticate-then-encrypt approach, but also encrypt-then-authenticate is possible.

The MAC is computed over the data(compressed or not), the TLS sequence number and the key used for the MAC computation. The padding is also part of the MAC computation to avoid those attacks that change the padding.

The MAC is then encrypted with the dedicated symmetric key and a suitable initialization vector(IV)

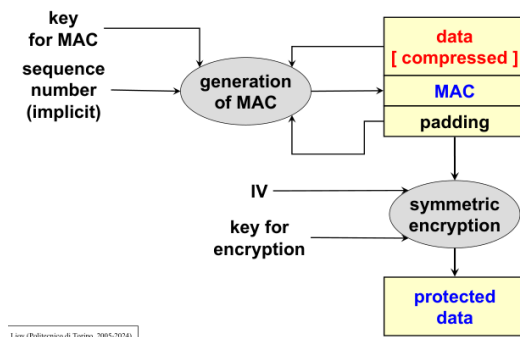


Figure 1.3: TLS data protection.

The keys are **directional**, so there are two keys( one for client to server and one for server to client) to protect against reuse of the sequence number in the opposite direction.

## 1.4 Relationship among keys and sessions

When a new session is created using the handshake protocol, a new **pre-master secret** is established using public key cryptography. Then from the session a new connection is created, which requires a random number to be generated, and exchanged between the client and the server. Those two values are combined via a KDF, usually HHKDF( HMAC-based key derivation function) to generate the **master secret**. This computation is done only once, and the secret is common to several connections.

The pre-master secret is then discarded, and the keys necessary for the MAC computation, encryption and, if necessary, IVs will be derived from the master secret.

You can notice that the master secret is common to any connections inside a session, but the per-connection keys are different every time. This is another important feature to avoid replay attacks, possible because numbering is per-connection. This solution also allows to reduce the cost of establishing new keys for each connection.




Figure 1.4: Relationship among keys and sessions.

## 1.5 Perfect Forward Secrecy

Since the keys are generated from symmetric crypto, if the private key used to perform encryption and decryption of the pre-master secret is compromised, all the previous communication can be decrypted because it is possible to derive the master-secret. This is only possible if the server has a certificate valid for both signature and encryption. In this context, perfect forward secrecy is desirable.

**Perfect Forward Secrecy** is a property of key-agreement protocols ensuring that the compromise of the secret key used for will compromise only current (and eventually future) traffic but not the past one

The most common way to achieve this is to use **ephemeral keys**, which are one-time asymmetric keys( used for key exchange). This means that the key pair used for key exchange is not a long term key pair, but a temporary one generated on-the-fly when necessary.

The ephemeral key needs to be authenticated, so only for this purpose the long-term key is used for signing the ephemeral key. This is done using DHKE instead of RSA because the latter one is really slow, while the former one is faster with the compromise of only using the established key for a certain number of session.

Let's now go over some considerations: if the temporary key is compromised, perfect forward secrecy of the communication is still valid because he can only decrypt the traffic exchanged using the temporary key. On the contrary, if the long-term key is compromised, no secret is really disclosed, because no traffic has been exchanged using it for encryption, but is still a problem for server authentication.

## 1.6 The protocol

The TLS handshake is always initiated by the client.


### 1.6.1 Client Hello and Server Hello

In version 1.2 the client sends a **Client Hello**, which contains:


- the SSL version preferred by the client, and the highest supported(2=SSL-2, 3.0=SSL-3, 3.1=TLS-1.0, ...)
- a 28 bytes pseudo-random number, which is the client random
- a session-id, which is  $0$  if the client is starting a new session, and  $1$  if the client is trying to resume a previous session
- a list of cipher suites supported by the client, in order to let the server choose the most secure one (the set of algorithms used for encryption, for key exchange, and for integrity)
- a list of compression methods supported by the client (supported only up to TLS 1.2)

And then a **server hello** is sent back, which contains:

- the SSL version chosen by the server, the highest one supported by both the client and the server
- a 28 bytes pseudo-random number, which is the server random
- a session identifier(session-id), which is a new one if the server is starting a new session, and the same as the client's if the server is resuming a previous session
- the cipher suite chosen by the server, the strongest common one between the client and the server
- the compression method chosen by the server



(a) The TLS handshake protocol(TLS 1.2).



(b) The TLS handshake protocol(TLS 1.3).

## 1.6.2 Cipher suite

A cipher suite is a string which contains the set of cryptographic algorithms used in the TLS protocol. A typical cipher suite consists of a key exchange algorithm, the symmetric encryption algorithm, and the hash function used for generating MACs. Some example of those are:

- SSL\_NULL\_WITH\_NULL\_NULL (no protection, used for the record protocol to be used in the handshake)
- SSL\_RSA\_WITH\_NULL\_SHA
- SSL\_RSA\_EXPORT\_WITH\_RC2\_CBC\_40\_MD5
- SSL\_RSA\_WITH\_3DES\_EDE\_CBC\_SHA

## 1.6.3 Certificates

After the initial exchange, the server is ready to authenticate itself.

The server sends its long-term public key certificate to the client for server authentication. Actually, the whole certificate chain, up to the root CA, must be sent. Furthermore, the subject of the certificate must match the server name.

Server authentication is implicit, because its private key is used to decrypt the pre-master secret, while client authentication is always explicit.



The implicit server authentication is based on the fact that the MAC is computed through the key derived through knowledge of the private key, so only the server can compute it.

Optionally, the server can request a certificate from the client for client authentication. In this case the server specifies the list of trusted CA's, and the client sends its certificate chain. The browsers show to the users (for a connection) only the certificates issued by trusted CAs. If client certificate verification is required, an explicit request to send the hash computed over all the handshake messages before this one and encrypted with the client private key is sent to the client.

#### 1.6.4 Key exchange

The key exchange is the most important part of the handshake protocol. If the server is using RSA for key exchange, the client generates a pre-master secret, encrypts it with the server's public key (which can be ephemeral or from its x.509 certificate) and sends it to the server. If RSA is not used, DHKE can be used to generate the pre-master secret, and in this case the server computes the value independently and the two parts can derive the master secret.

Another option is to use FORTEZZA, which is a key exchange algorithm based on DH.

#### 1.6.5 Certificate verify

In case the server requested client authentication, the client will be required to send the certificate to prove that he is the owner of its private key. The message to be signed is the hash of all the messages exchanged up to this point in the handshake protocol. This is to avoid replay attacks too.

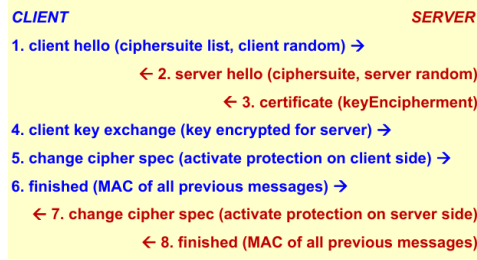
#### 1.6.6 Change cipher spec

The change cipher spec message is used to trigger the change of the algorithms to be used for message protection. It allows to pass from the previous unprotected messages to the protection of the next messages with algorithms and keys just negotiated, thus is technically a protocol on its own and not part of the handshake. Some analysis even say that it could be removed from it.

#### 1.6.7 Finished message

The finished message is the last message of the handshake protocol, and the first message protected by the negotiated keys and algorithms. It is necessary to ensure that the handshake has not been tampered with, and it contains a MAC computed over all the previous handshake messages (but change cipher spec) using as a key the master secret. Notice that the finished message is different for the client and the server, because the MAC is computed over different messages.

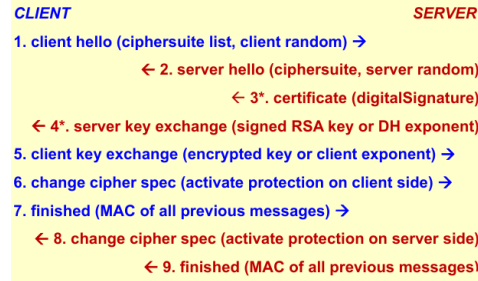
This allows to prevent rollback man-in-the-middle attacks (version downgrade or cipher-suite downgrade)



(a) TLS handshake(no ephemeral, no client authN).



(b) TLS handshake(no ephemeral, client authN).




(c) TLS handshake(ephemeral, no client authN).


Figure 1.6: TLS handshake protocol.

## 1.7 Setup Time

The setup time is the time required to establish a secure connection between the client and the server. TLS depends on TCP, so the TCP handshake must be taken into account. Then the TLS handshake is performed, meaning that typically 3 RTTs (1 for TCP and 2 for TLS) are required to establish a secure connection. Usually after 180ms the two parties are ready to send protected data( assuming 30ms delay one-way).



(a) TLS link teardown.



(b) TLS resume session.

Figure 1.7: TLS link teardown and resume session.

## 1.8 TLS versions

### 1.8.1 TLS 1.0

TLS 1.0, or SSL 3.1, was released in 1999. It is the first version of the protocol, and it is based on SSL 3.0. Previous version were using proprietary solutions, so the adoption of open standards was strongly encouraged.

### 1.8.2 TLS 1.1

TLS 1.1 was released in 2006, and it introduced some security fixes especially to protect against CBC attacks. In fact, the implicit IV is replaced with an explicit IV to protect against CBC attacks. Also protection against padding oracle attacks were introduced to reduce the information leaks. For this reason Passing errors now use the `bad_record_mac` alert message (rather than the `decryption_failed` one). Furthermore, premature closes no longer cause a session to be non-resumable.

### 1.8.3 TLS 1.2

TLS 1.2 was released in 2008, and it introduced some new features and improvements. The ciphersuite also specifies the pseudo random function instead of leaving the choice to the implementation. The sha-1 algorithm was replaced with SHA-256, and it also added support for authenticated encryption, such as AES in GCM or CCM mode.

All the ciphersuites that use IDEA and DES are deprecated.

## 1.9 TLS attacks

### 1.9.1 Heartbleed

Heartbleed is a security bug in the OpenSSL cryptography library, which is a widely used implementation of the TLS protocol. It was able to exploit the fact that the heartbeat extension keeps the connection alive without the need to negotiate the SSL session again. The attacker could send a heartbeat request, but the length of the response is much longer (up to 64KB) than the actual data sent by the client. This attack could then allow to leak memory contents.

### 1.9.2 Bleichenbacher attack

Blackenbacher is a 1998 attack and it's the so-called the million message attack because it exploited a vulnerability in the way the RSA encryption was done. RSA requires the padding to be done in a certain way, because if it is unoptimally done, it could cause some issues.

The attacker could perform an RSA private key operation with a server's private key by sending a million or so well-crafted messages and looking for differences in the error codes returned. By basically knowing the public key and trying to decrypt a message with some guessed private keys, the different responses obtained were giving hints about which bits were correct and which bits were wrong.

Later on the RSA implementations moved to RSA-OAEP, which is a padding scheme that is provably secure against chosen-ciphertext attacks.

In 2017 another variant of this attack was discovered, called ROBOT (Return Of Bleichenbacher's Oracle Threat), to which many major websites, like Facebook, were vulnerable.

### 1.9.3 Other attacks against SSL/TLS

Some other attacks against SSL/TLS are CRIME, BREACH, BEAST and POODLE.

**Crime** is an attack against the **compression algorithm** used in **SSL/TLS**, which by injection chosen plaintext in the user requests, for example by using a form or choosing fraudulently an username that is displayed, and then measure the size of the encrypted traffic, an attacker could recover specific plaintext parts exploiting information leaked from the compression, and this is part of the reason why the compression is deprecated in TLS 1.3.

**BREACH** is an attack against the **HTTP compression** to deduce a secret within the HTTP response provided by the server. It is different from Crime because the former is an attack against the compression algorithm used in SSL/TLS, while the latter is an attack against the HTTP compression.

**BEAST** is an attack that exploits a vulnerability in the way the **CBC** mode of operation is used in SSL/TLS. The attack is possible if **IV concatenation** is used, meaning that the initial vector for the next encryption is taken from the end of the previous encryption. A MITM may decrypt HTTP headers with a blockwise-adaptive chosen-plaintext attack, and by doing so, he's able to decrypt HTTPS requests and steal information such as session cookies.

**POODLE**, or Padding Oracle On Downgraded Legacy Encryption, is an attack that exploits the fact that SSL 3.0 uses a padding scheme that is vulnerable to a **padding oracle attack**, by acting as a MITM. This is done by exploiting SSL-3 fallbacks to decrypt data. This is also the only attack among those that still works today.

**FREAK**, or Factoring RSA Export Keys, is an attack that exploits the downgrades on TLS to export-level RSA keys to a factorizable bit length(512 bits). It is also possible to carry this out by downgrading the symmetric key too and then perform a brute force attack(40-bit). As you can see from figure 1.8, in the first phase the random and the supported elliptic curve are not altered by the MITM, but only the supported cipher suites are altered( to export level ones). Usually 40 bits cipher suits should not be configured at all, but some misconfigurations may happen. The third phase is where the magic happen: we have theoretically the MAC in the FINISHED message to protect against tampering (recall that the MAC will be computed over all the handshake messages) but since the MAC is protected with the master secret, the master secret is only 40 bits. The attacker can then brute force the master secret on-the-fly, recompute the MAC so that the server will accept the message. Since the attacker have access to the premaster secret and the master secret, all traffic beyond this point is encrypted with a weak shared key and the middleman can read and even modify the traffic.

For all those reasons SSL-3 has been disabled on most browsers, but its still needed for some browsers, for example IE6 by Microsoft, which is outlasting its expected life span, because its the default browser on Windows XP, which is still used today unfortunately, meaning that the window of exposure for those attacks is still open.

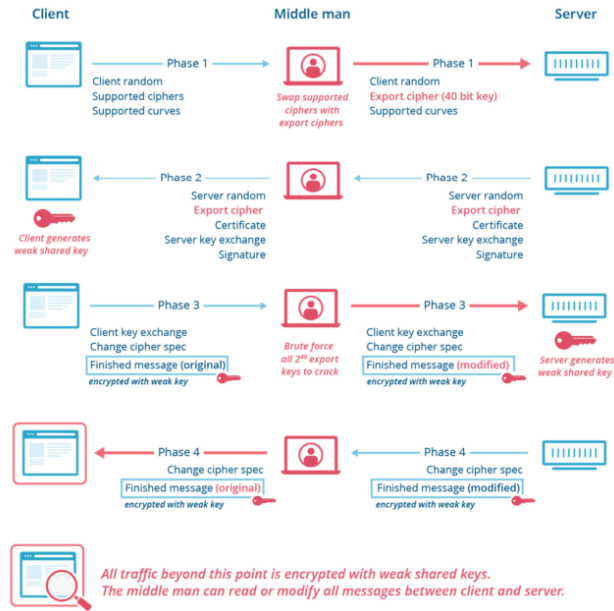


Figure 1.8: FREAK attack.

## 1.10 ALPN extension

The ALPN extension, or Application-Layer Protocol Negotiation, is an extension that allow to negotiate the application protocol to speed up the connection creation, avoiding additional round-trips for application negotiation. It is used to negotiate the protocol to be used on top of the TLS connection, such as HTTP/2, SPDY, or QUIC, before the connection is established. This is useful because it saves time, obviously, because after the connection is established, the client and server could still fail to communicate because the application protocol is not supported by the server.

The extension is inserted in the client hello message, by setting the ALPN flag to true and providing a list of supported protocols in the client HELLO message. The server will respond with the ALPN flag set to true(if it supports the extension) and the selected protocol.

This is also useful for those servers that use different certificates for the different application protocols.

## 1.11 TLS False Start

TLS False Start is another extension that allows the client can send application data together with the ChangeCipherSpec and Finished messages, in a single segment, without waiting for the corresponding server messages. The biggest advantage of using this is the reduction of the latency to 1 RTT. In theory this should work without changes, but to use this in Chrome and Firefox they require the ALPN and the Forward Secrecy enabled, while Safari requires forward secrecy.

## 1.12 The TLS downgrade problem

In theory, when negotiating the TLS version to be used, the client sends (in ClientHello) the highest supported version, while the server notifies (in ServerHello) the version to be used (highest in common with client).

For example if the client support up to SSL 3.3(TLS 1.2) and the server support the same version, the connection will be established using TLS 1.2. But if the server supports only SSL 3.2(TLS 1.1), the connection will be established using SSL TLS 1.1.

Some servers, instead of sending the highest version supported, just close the connection, forcing the client to retry with a lower version of the protocol. An attacker could exploit this behavior to force the client to use an older version of the protocol, by repeatedly closing the connection, and then exploit the vulnerabilities of the older version of the protocol.

This means that its not a problem of the protocol itself, but of the implementation of the server.

### 1.12.1 TLS Fallback Signalling Cipher Suite Value (SCSV)

This behavior is not always an attack, for example there could be an error in the channel, which is closed by the server. This means that there's a need to distinguish between a real attack and a simple error.

The **TLS Fallback SCSV** is a **special value** that is used to prevent the protocol downgrade attacks, not only cipher suite. It does so by sending a new (dummy) cipher suite value(TLS\_FALLBACK\_SCSV) which is sent by the client when opening a downgraded connection as the last value in the cipher suite list.

If the server receives this value and still supports a higher version of the protocol, it will know that the client is trying to downgrade the connection, and it will refuse to establish the connection by sending an **inappropriate fallback** alert message and closing the channel. This notifies the client that he should retry with the highest version of the protocol supported by himself.

Many servers do not support SCSV yet, but most servers have fixed their behavior when the client requests a version higher than the supported one so browsers can now disable insecure downgrade

## 1.13 TLS session tickets

We know that session resumption is possible with TLS, but the server needs to keep a cache of session IDs, which may become very large for high traffic servers. For this reason, the **TLS session tickets** were introduced, which are an extension allowing the server to send the session data to the client encrypted with a server secret key. This data is stored by the client, which will send it again when it wants to resume a session. This allows to move the cache to the client side. Obviously, this data has to be encrypted with a server secret key. Even if this behaviour is desirable for the server, it still needs to be supported by the browser (it's an extension after all) and needs a mechanism to share keys among the servers in a load-balancing heavy environment.

## 1.14 The Virtual Server Problem

Nowadays, virtual servers are very common in web hosting, because they allow to have different logical names associated with the same IP address( ie: home.myweb.it=10.1.2.3, food.myweb.it=10.1.2.3). This is easy to manage in HTTP/1.1 but quite troublesome with HTTPS, because TLS is activated before the HTTP request is sent, which makes it difficult to know which certificate should be provided in advance. The solutions are quite simple:

- use a wildcard certificate, which is a certificate that is valid for all the subdomains of a domain (ie: \*.myweb.it)
- use the SNI (Server Name Indication) extension, which is an extension that allows the client to specify the hostname of the server it is trying to connect to, allowing the server to provide the correct certificate. This is sent in the ClientHello message.
- provide a certificate with a list of servers in subjectAltName, which allow to share the same private key for different servers.

## 1.15 TLS 1.3

TLS 1.3 was released in 2018, and it introduced some new features while solving some of the most common problems of the previous versions:

- **reduce the handshake latency** in general
- **encrypting** more of the **handshake** (for security and privacy, after all up to TLS 1.2 the handshake was in clear text)
- **improving resiliency** to cross-protocol attacks
- removing legacy features

We will now go over the main changes in TLS 1.3.

### 1.15.1 Key exchange

In this version, the support for static RSA and DH key exchange was removed for many reasons:

- it does not implement forward secrecy
- its difficult to implement correctly, which is a problem because it exposes the system to many attacks like Bleichenbacher

Now Diffie-Hellman ephemeral (DHE) and Elliptic Curve Diffie-Hellman Ephemeral (ECDHE) are the only key exchange methods supported, with some required parameters( some implementations weren't really up to the standards, ie: used DH with small numbers( just 512 bits) or generated values without the required mathematical properties).

### 1.15.2 Message protection

TLS 1.3 greatly improves message protection by eliminating several vulnerabilities present in earlier versions. Previously, issues arose from using CBC mode with authenticate-then-encrypt, which led to attacks like Lucky13 and POODLE. The use of RC4 also allowed plaintext recovery due to measurable biases, and compression enabled the CRIME attack. TLS 1.3 addresses these weaknesses by removing CBC mode and enforcing AEAD modes for stronger security. Insecure algorithms such as RC4, 3DES, Camellia, MD5, and SHA-1 have been dropped, and compression is no longer used, ensuring a more secure cryptographic environment.

### 1.15.3 Digital signature

In TLS 1.3, digital signatures have been strengthened to address earlier vulnerabilities. Previously, RSA signatures were used on ephemeral keys with the outdated PKCS#1v1.5 schema, leading to potential flaws. The handshake was authenticated using a MAC instead of a proper digital signature, which exposed the protocol to attacks like FREAK.

TLS 1.3 improves this by using the modern, secure RSA-PSS signature scheme. Additionally, the entire handshake is signed, not just the ephemeral keys, providing more comprehensive security. The protocol also adopts modern signature schemes, enhancing the overall strength of the cryptographic process.

### 1.15.4 Ciphersuites

TLS 1.3 simplifies the protocol by reducing the complexity seen in earlier versions, which had a long list of cryptographic options that grew exponentially with each new algorithm. This complexity made configuration and security management difficult.

To address this, TLS 1.3 specifies only essential, orthogonal elements: a cipher (and mode) combined with an HKDF hash function. It no longer ties the protocol to specific certificate types (like RSA, ECDSA, or EdDSA) or key exchange methods (such as DHE, ECDHE, or PSK).

Moreover, TLS 1.3 narrows the selection to just five ciphersuites:

- TLS\_AES\_128\_GCM\_SHA256
- TLS\_AES\_256\_GCM\_SHA384
- TLS\_CHACHA20\_POLY1305\_SHA256
- TLS\_AES\_128\_CCM\_SHA256
- TLS\_AES\_128\_CCM\_8\_SHA256 (deprecated but yet supported for computational environments with low capacity)

### 1.15.5 EdDSA

**EdDSA**, or Edwards-curve Digital Signature Algorithm, is a digital signature scheme using the EdDSA signature scheme, which is a variant of the standard elliptic curve DSA schema. The EdDSA scheme, unlike standard DSA, doesn't require a PRNG, which could in some



cases leak the private key if the underlying generation algorithm is broken or predictable. EdDSA picks a nonce based on a **hash of the private key** and the **message**, which means after the private key is generated there's no need anymore for random number generators. Another advantage is that the EdDSA is faster in signature generation and verification than the standard DSA, because it implements simplified point addition and doubling.

EdDSA is using an n-bit private and public keys and will generate a signature which has a size which is the double of the keys size. As per the RFC stipulations, EdDSA must support the Ed25519 and Ed448 curves, which are based on the Edwards-curve Digital Signature, which uses respectively SHA-2 and SHA-3 as hash functions. Among those two, Ed25519 is the most used, which is also used for an implementation of EcDH.

When it comes to the RFCs there's an issue: there are two RFCs that are slightly different, one is the RFC 8032, which is for general internet implementations, with the implementation details left to the developers, and there's also FIPS 186-5, which specifies not only the mathematics behind an algorithm but also implementation details because it is a standard for the US government.

### 1.15.6 Other improvements

TLS 1.3 introduces several key improvements to enhance security and efficiency. One major change is that all handshake messages sent after the ServerHello are now encrypted, ensuring better confidentiality during the handshake process, even if the keys have not been established (we will go over the details later). The new **EncryptedExtensions** message allows additional extensions, which were previously sent in cleartext during ServerHello, to also benefit from encryption, improving privacy.

Additionally, key derivation functions have been redesigned to support easier cryptographic analysis due to their clear key separation properties. HKDF is now used as the underlying primitive for key derivation, providing a robust and flexible method.

The handshake state machine has also been overhauled, making it more consistent by removing unnecessary messages, such as **ChangeCipherSpec**, which is now only retained for compatibility with older systems. This restructuring streamlines the handshake process, reducing complexity and improving protocol efficiency.

### 1.15.7 HKDF in TLS 1.3

HKDF is one of the novelties of TLS 1.3, and is a HMAC-based extract-and-expand Key Derivation Function. As a normal key derivation function, it takes 4 inputs: the salt, the input keying material (which is basically a secret key), the info string and the output length. It basically takes the first 2 inputs to derive a key which will be used to derive a pseudorandom key. Then the second stage expands this key into several additional pseudorandom keys, which are the output of the KDF. So multiple outputs can be generated from the single IKM value by using different values info strings.

By calling the function repeatedly call the HMAC function with PR keys and the info string as the input message, and then concatenate the results, by appending the output of the HMAC function to the previous output and incrementing an 8-bit counter.

This means that from one call we can generate at most  $256(2^8)$  different keys.

The finished message is not protected with the master secret as it was in TLS 1.2, but with

the specific key, which is created with expand starting from the base key and putting as info the text "finished" and then the desired length. The pre-shared key to protect the ticket contains the word "resumption". So you start with some basic key material and then by using different labels, you are able to generate the different keys.

```
HKDF-Expand( HKDF-Extract ( salt, IKM ), info, length )
```

Listing 1: HKDF-Expand pseudocode.

### 1.15.8 TLS-1.3 handshake

Even the handshake in TLS 1.3 was changed. The first thing that we notice is that the change cipher specs has been removed from the standard.

At first, as you can see from figure 1.9, the connection starts with an hello message from the client, which will also send it's list of ciphersuites and the key share material, in which the client will send it's DH parameters(it is still needed after all). All this data is sent in the same message to save some RTTs and to allow the middleboxes with lower TLS version support to understand the message, interpreting them as an extension of the client hello and thus skipping them. In fact most of the TLS 1.3 features are implemented in message extensions.

The following messages are just the specular ones but sent from the server this time. After this point it is possible to have confidentiality of the communication, because a shared key has been established(the curly braces in the picture are just for that).

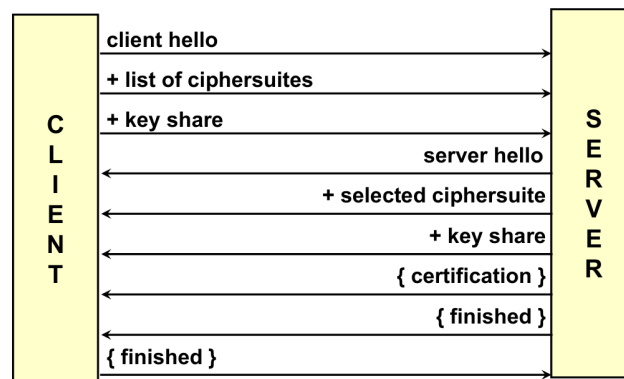


Figure 1.9: TLS 1.3 handshake.

#### Client request

The client hello message contains the **client random**, the highest supported protocol version(which is 1.2), client random, the supported cipher suites and compression methods(no compression allows but still necessary for backward compatibility) and of course the session id.

The supported extensions are:

- key\_share = client (EC)DHE share

- `signature_algorithms` = list of supported algorithms
- `psk_key_exchange_modes` = list of supported modes for the pre-shared key exchange
- `pre_shared_key` = list of PSKs offered, which are not the real keys, are just an index label for the pre-shared keys available

Recall that if the extension is not supported by the server or middlebox it will be ignored.

### Server response

The server will respond with its parameters, some of which will already be encrypted because the key has been established. The key exchange material will be sent in clear of course (the `key_share` and the `pre_shared_key`).

The server's parameters are present too and encrypted, which are

- `{ EncryptedExtensions }` = responses to non-crypto client extensions
- `{ CertificateRequest }` = request for client certificate, optional and encrypted for privacy

Even the server authentication parameters are sent in the same message:

- `{ Certificate }` = X.509 certificate (or raw key, RFC-7250, useful for low power IoT devices)
- `{ CertificateVerify }` = signature over the entire handshake
- `{ Finished }` = MAC over the entire handshake

Finally, the server can immediately send application data, if needed, which will be encrypted with the key derived for data protection, not the one derived for the handshake.

### Client finish

The client will respond with the last messages, which are:

- `{ Certificate }` = X.509 certificate (or raw key, RFC-7250)
- `{ CertificateVerify }` = signature over the entire handshake
- `{ Finished }` = MAC over the entire handshake

### Pre-shared keys

In TLS, the pre-shared key (PSK) replaces the session ID and session ticket. PSKs are agreed upon during a full handshake and can be reused for multiple connections.

Pre-shared keys are only used if a previous session has been established.

They can be combined with (EC)DHE to achieve forward secrecy, where the PSK is used for authentication and (EC)DHE for key agreement. While PSKs can be generated out-of-band (OOB) from a passphrase, this is risky due to the potential lack of randomness, making brute-force attacks feasible. Therefore, using OOB PSKs is generally discouraged.

## 0-RTT connections

In TLS 1.3, when using a pre-shared key (PSK), a client can send "early data" with its initial message, which is protected by a specific key. However, this approach lacks forward secrecy because it relies solely on the PSK and could be vulnerable to replay attacks. While some complex mitigations exist, they are particularly challenging for multi-instance servers.

## Incorrect share

In TLS 1.3, if a client sends a list of (EC)DHE groups that the server does not support, the server responds with a HelloRetryRequest, prompting the client to restart the handshake with different groups. If the new groups are also unacceptable, the handshake will be aborted, and the server will send an appropriate alert.

## 1.16 TLS and PKI

Even with TLS 1.3, public key certificates are required, meaning TLS has a connection with Public Key Infrastructure (PKI), as certificates are typically issued by a PKI. PKI is always needed for server authentication and optionally for client authentication unless using "true" PSK (Pre-Shared Key) authentication. In this case, "true" refers to manually pre-installed shared keys between the client and server, which is rare, often limited to certain embedded systems.

When a peer (server or client) sends its certificate to the other peer:

- The entire certificate chain is sent, which includes the peer's certificate and certificates of intermediate CAs (Certification Authorities). However, the root CA certificate is not sent, as it should already be trusted and present on the other side. Some attackers may attempt to send a fake root CA certificate; if the receiving peer accepts it without proper verification, the attack is successful. For instance, in TLS monitoring systems, it's crucial to ensure that the certificate chain never includes the root CA.
- The receiving peer must validate each certificate in the chain. A potential vulnerability is when some browsers and servers only validate the peer's end certificate, neglecting the intermediate CAs, which could be invalid. Proper validation must recursively verify each certificate in the chain, including checking the correctness of signatures, issue, and expiration dates.
- It's also essential to verify whether a seemingly valid certificate has been revoked. This can be done using either:
  - **CRL (Certificate Revocation List):** A list of revoked certificates, though the list can be large, making downloading and storing CRLs cumbersome—especially since a separate CRL is needed for each step in the chain.
  - **OCSP (Online Certificate Status Protocol):** More efficient performance-wise, as it allows querying an OCSP server to check if a certificate is valid. However, this method raises privacy concerns, as the OCSP server gains visibility into the servers being visited.

Both CRL and OCSP add additional network requests and delays to the setup, as the certificates must be validated before the session begins. For example, OCSP can introduce a median delay of +300ms, with an average of +1 second. Therefore, privacy and performance trade-offs must be considered.

## 1.17 TLS and certificate status

What should be done if the CRL or OCSP servers are unreachable? This could happen due to server or network failures, or even because a firewall is blocking access—common in companies with strict firewall rules due to security policies. Another reason could be that OCSP responses are sent in clear text, even though they are signed. The potential responses to this situation are:

- **Hard fail:** The page is not displayed, and the user is shown a security warning. For instance, "Sorry, we cannot display this page because we couldn't verify its revocation status." While this is the most secure option, it is also the most inconvenient for users. It could also result from a (D)DoS attack on the OCSP or CRL server.
- **Soft fail:** The page is shown regardless of whether the certificate's revocation status can be verified, assuming the certificate is still valid.

Both approaches lead to extra load time, often longer than just waiting for an OCSP response. This happens because the system needs to wait for the timeout after initiating a TCP connection. To address these issues, two different solutions are typically used.

### 1.17.1 Pushed CRL

In most cases, revoked certificates result from a compromised intermediate CA. When this happens, all certificates issued by that CA must be revoked, which can be a significant number. As a result, browser vendors have chosen to push updates containing only major revoked certificates. If an intermediate CA becomes invalid, it is promptly marked as invalid, and the CRL containing this information is distributed through a browser update. Here are some examples of how different browsers handle this:

- **Internet Explorer:** Updates its revoked certificate list with each browser update, which isn't ideal. If an update is delayed, the user may not be informed about the CA revocation in a timely manner.
- **Firefox:** Uses a system called *oneCRL*, which is part of its blocklist management process. This list, maintained by Firefox's developers, is updated nightly and includes CRLs of revoked intermediate CAs as well as known malicious websites.
- **Chrome:** Utilizes *CRLsets* to manage these CRLs, which are pushed with each browser update. On startup, Chrome checks for new CRLsets.

## 1.18 OSCP Stapling

Let's now look at how browsers and servers manage OCSP-related issues, which can affect both verification and privacy. The term "Stapling" refers to keeping things together.

- Automatic downloads of CRL and OCSP are often disabled because they take too long. As a result, browsers may be vulnerable to attacks using revoked certificates.
- Pushed CRLs only include a subset of revoked certificates; otherwise, the list would become too large.
- Browser behavior varies significantly in this area. Depending on the browser, different features may be available, and there is no standardized method for handling certificate revocation.

To improve this, instead of relying on the client to fetch revocation information, the server takes responsibility through a process called *OCSP stapling*. When the server sends its certificate, it also provides recent revocation information via an OCSP response (essentially saying, "Here's my certificate, and here's proof it's valid"). This allows the client to skip making a separate OCSP request, as the necessary information is included in the initial TLS handshake.

OCSP Stapling is a TLS extension, meaning it is not part of the standard handshake and must be supported by both the server and the client. It is announced during the TLS handshake when the server sends the *Server Hello* message, indicating to the client that it will be using this extension.

The first version of OCSP Stapling is defined in RFC-6066 (extension **status\_request**), while version two is in RFC-6961 (extension **status\_request\_v2**) with the value **CertificateStatusRequest**. The TLS server pre-fetches the OCSP response from an OCSP server and includes it in the handshake as part of the server's certificate message. This message contains not only the certificate chain but also the OCSP responses for each certificate in the chain. As a result, the message size increases since there must be an OCSP response for each certificate, which validates that the certificates have not been revoked. This process "staples" the OCSP responses to the certificates.

The main advantage is improved privacy for the client, as it no longer needs to request information directly from the OCSP server—this is handled by the TLS server, meaning the OCSP server remains unaware of which clients are requesting access. Additionally, the client avoids having to establish a separate connection to the OCSP server, improving both speed and privacy.

However, a downside is that the freshness of the OCSP responses can be an issue. Since the OCSP response is pre-fetched by the server, it may not always be up-to-date (the server might store the response for a few minutes). This delay creates a small window of opportunity for attacks. For example, an attacker could retrieve the valid certificate and OCSP response, then attack the server to steal the private key and set up a fake server. When users connect to this fake server, the attacker can still present a valid certificate and OCSP response until the browser accepts it (would a browser accept an OCSP response from 30 minutes ago?).

OCSP stapling is typically handled automatically by the server, though the client can explicitly request it. When the client connects, it may not know whether the server supports OCSP stapling, so the client can send a *Certificate Status Request* (CSR) as part of the *ClientHello* message to request the OCSP responses in the TLS handshake.

This process is optional: even if the client requests OCSF responses, the server may not support stapling. In this case, the server might return the OCSF responses for its certificate chain in a separate message called *CertificateStatus*.

The issues include:

- Servers may ignore the status request; they are not obligated to provide a response.
- Clients may proceed with the handshake even if no OCSF response is provided.

Although the technology exists, there is uncertainty about whether it is widely implemented. It is important to monitor the process to verify whether OCSF responses are being provided, requested, or ignored, as this offers insight into the actual security status, not just the theoretical one. To address this uncertainty, a newer solution called *OCSF Must Staple* has been introduced, which forces OCSF stapling.

### 1.18.1 OCSF Must Staple

When a server requests a certificate from a certification authority, it may declare that it will always provide OCSF responses to clients, even if not explicitly requested. This information can be included in the certificate. Such servers present a certificate to the client that contains a certificate extension called **TLSFeatures**, as defined in RFC-7633. This extension informs the client that every time it connects to that server, it must receive a valid OCSF response as part of the TLS handshake. If the client does not receive an OCSF response, it should reject the server's certificate. Essentially, the server promises to always send both its certificate and the OCSF response. If the server does not provide the OCSF response, it can be considered a fraudulent server.

The benefit of this approach is that the client no longer needs to query the OCSF responder, and it enhances security by preventing attacks that block OCSF responses for a specific client or DoS attacks against the OCSF responder.

### Key actors and their responsibilities

- **Certification Authority (CA):** Must include the **TLSFeatures** extension in the server certificate if requested by the server's owner.
- **OCSF Responder:** Must be available 24/7 to provide valid OCSF responses. This requires a robust infrastructure with multiple network access points, backup servers, and resilience against DoS attacks. Otherwise, servers would not be able to establish connections, as they must provide OCSF responses.
- **TLS Client:** Must send the Certificate Status Request (CSR) extension in the *ClientHello* message, recognize the OCSF Must-Staple extension (if present in the server's certificate), and reject the server's certificate if no OCSF response is provided when promised.
- **TLS Server:** Must pre-fetch and cache OCSF responses for a certain period, and include an OCSF response in the TLS handshake. It must also handle errors when communicating with OCSF responders, in case the responder is temporarily unavailable.

- **Server Administrators:** Should configure their servers to use OCSP Stapling and request a server certificate that includes the OCSP Must-Staple extension.

One potential issue is the duration of the OCSP stapled response (for example, Cloudflare has a validity of 7 days). A window of 7 days may be problematic if a server is compromised, as there is potential exposure during that time.

## 1.19 Questions and answers

### Question 1

what of the following remarks regarding TLS are possibly true?

- ✓ has been established a session S1 that involved connection C1 and C2 at time t1
- ✓ has been established a session S2 that involved connection C1, C2 and C3, at different and non- overlapping time intervals
- × connection C5 has been used initially in session S3 and then re-used in session S4 for performance-sake
- × in connection C1.1 of session S1 has been used AES192, while in Connection C1.2 of the same session S1 has been used AES128 since regarded less sensible information

### Question 2

focusing on TLS Connections and Sessions

- ✓ sessions are typically relatively long-life in respect to connections
- × connections allow to re-use cryptographic parameters defined during handshake, thus reducing significantly the handshake phase
- × by using resumption mechanisms, the client can resume a connections decreasing the overall connection time
- ✓ each connection has its own specific set of parameters like sequence numbers and keys for integrity and confidentiality

### Question 3

focusing on Perfect Forward Secrecy

- ✓ starting from version TLS version 1.3, it must be always enabled
- × using RSA mechanisms in TLS, it would be theoretically impossible to achieve
- ✓ using RSA mechanisms in TLS, it would be impractical to achieve
- × using ECDH mechanisms in TLS, it would be impractical due to the length of the key parameters



#### Question 4

Which of the following properties is NOT directly related to Perfect Forward Secrecy in a TLS context?

- × the use of ephemeral keys
- ✓ protection against replay attacks
- × security of past sessions if the server's private key is compromised
- × independent session keys for each connection

#### Question 5

In the context of TLS 1.3, what is the role of PFS in session resumption mechanisms like O-RTT?

- × O-RTT resumption is vulnerable to replay attacks, thus does not support PFS
- × O-RTT session resumption uses pre-shared keys that provide Perfect Forward Secrecy.
- ✓ O-RTT session resumption compromises Perfect Forward Secrecy for faster handshakes.
- × O-RTT session resumption reuses the original session's ephemeral keys

#### Question 6

Which TLS feature was introduced to specifically mitigate downgrade attacks?

- × HSTS (HTTP Strict Transport Security)
- × OCSP Stapling
- × Certificate Pinning
- ✓ TLS FALLBACK SCSV

#### Question 7

Given the following packet capture:

- × indicate a
- × might be the initial phase of a secured real-time data exchange (like video streaming)
- × It refers to an aborted TLS session, due to the impossibility to verify the x509v3 certificate
- ✓ It refers to an aborted TLS session, due to TLS version mismatch between the versions supported by the server and by the client

# Chapter 2

## SSH

SSH, or Secure Shell, is a **network protocol** that allows for **secure communication** between two computers, which makes it a "*competitor*" of the much more used TLS protocol. It has been introduced in 1995 by Tatu Ylönen, a Finnish computer science student, as a response to a security incident per which the FTP credentials of his university were sniffed, but it has been commercialized since so it was not open source anymore. A open source fork called OpenSSH has been introduced in OpenBSD in 1999, and has been standardized in 2006 by the IETF in RFC 4251, referred to it as SSH-2.

As it is the most used version, here we will always refer to SSH-2 (unless otherwise noted)

### 2.1 Architecture

From an architectural point of view, SSH is a three layer architecture, which means it is simpler than TLS. SSH uses TCP as the Transport Layer Protocol, but the SSH transport layer provides a nice set of features, which are **initial connection**, **server authentication**, **confidentiality** and **integrity** and **key re-exchange** (RFC-4253 recommends after 1GB of data transmitted or after 1 hour of transmission, notice that this does not refer specifically to SSH but to any encrypted channel) which is crucial for long running connections.

Beware that the SSH transport layer is not the same as the TCP transport layer, but it is a layer that is built on top of the TCP transport layer, even though they have similar names.

There's also a specific protocol for **user authentication**, compulsory with SSH, and a **connection protocol** too, which supports multiple connections (channels) over a single secure channel (implemented with the transport layer protocol). This means that SSH can act as a site-to-site VPN.

#### 2.1.1 SSH Transport Layer Protocol

The general schema of an SSH connection is shown in figure 2.1.

The connection is setup with TCP, by default on port 22, on which the client initiates the

connection.

Then the SSH version string is exchanged by both sides, which contains the security features supported. Both sides must send a **version string** of the following form (the pipe symbol is used to separate the fields):

SSH-protoversion-softwareversion|SP|comments|CR|LF

SP is a space character, CR is a carriage return character and LF is a line feed character. If the comments are omitted, the SP is also omitted.

It is also used to trigger compatibility extensions and to advertise the protocol implementation capabilities.

A **key exchange** follows, which includes algorithm negotiation for the key exchange itself. After this point data can be exchanged, and the connection is closed when the data exchange is finished.

The termination of the TCP connection is not part of the protocol itself, but it's performed by TCP itself.




Figure 2.1: SSH Transport Layer Protocol

### Binary Packet Protocol

All packets that follow the version string exchange are sent using the Binary Packet Protocol, which is roughly the basic unit of transmission. Every message in this kind of packet is sent as a sequence of bytes. The Binary Packet Protocol is pretty simple, it consists of:

- the **length** of the **packet** (4 bytes), which does not include the MAC and the packet length field itself
- the **padding length** (1 byte), which is the length of the random padding field
- the **payload**, which may be compressed. Its size is the length of the packet minus the length of the length field minus 1 (for the padding length field), up to a maximum uncompressed size of 32768 bytes (32KB)




Figure 2.2: Binary Packet Protocol

- the **random padding**, minimum 4 bytes, up to the entire size of one block, and random to add confusion to the packet
- the **MAC**, computed over the clear text packet and the implicit sequence number

The first 4 fields are encrypted.

### Key exchange

Another essential part of the SSH protocol is the key exchange, which is used to establish the keys and to negotiate the algorithms to be used.

One peculiarity of the SSH protocol is that the algorithms used are selected based on directionality, which means that it is possible to use different algorithms for the client-to-server and server-to-client communication.

SSH-2 allows to select the key exchange function as well as the hash algorithm to be used in the key derivation function.

This means that the two parties have *"full"* control over the algorithms to be used (if supported).

Another interesting aspect of the protocol is that, as it was proposed as an alternative to the Telnet protocol, it included negotiation of the language to be used, because of the text-only nature of the protocol.

The last field of the key exchange (first key exchange packet follows) contains an attempt to guess the agreed key exchange algorithm.

It contains the following fields:

- SSH\_MSG\_KEXINIT
- cookie (16 random bytes)
- kex\_algorithms
  - eg: diffie-hellman-group1-sha1, ecdh-sha2-OID\_of\_curve

- `server_host_key_algorithms`: used for the server authentication
  - eg:ssh-rsa, ssh-dss, ecdsa-sha2-NISTP256
- `encryption_algorithms_client_to_server, ... server_to_client`
  - aes-128-cbc, aes-256-ctr, aead\_aes\_128\_gcm
- `mac_algorithms_client_to_server, ... server_to_client`
  - hmac-sha1, hmac-sha2-256, aead\_aes\_128\_gcm
- `compression_algorithms_client_to_server, ... server_to_client`
  - none, zlib
- `languages_client_to_server, ... server_to_client`
- `first_kex_packet_follows` (flag)

Keep in mind that because of the negotiation of the protocols to be used, SSH is potentially vulnerable to downgrade attacks.

A complete list of the supported algorithms can be found [here](#)

### Diffie-Hellman key agreement

This is also an **implicit authentication** of the server, due to the fact that the server can compute the correct premaster secret using the private key.

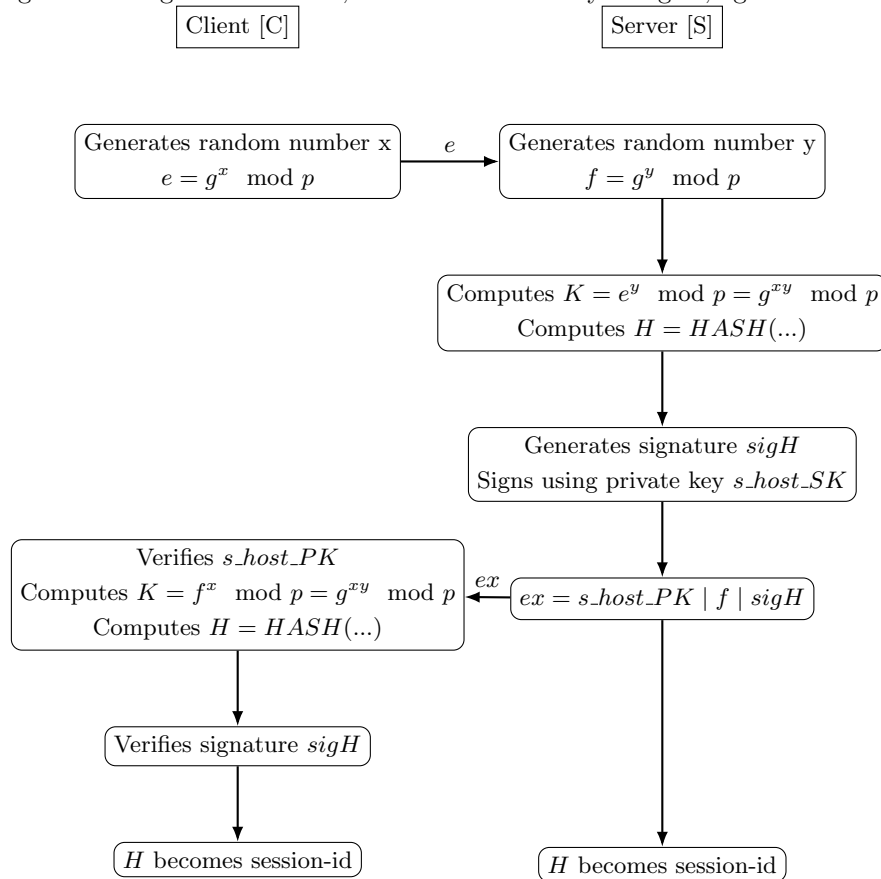
In this phase the client generates a random number  $x$  and computes  $e = g^x \bmod p$  and sends it to the server. The server generates a random number  $y$  and computes  $f = g^y \bmod p$  and sends it to the client. At this point the key  $K$  has been established. The server computes the **exchange hash**  $H$  over many fields(which makes it basically a keyed digest):

- the client version string
- the server version string
- the client's KEXINIT message
- the server's KEXINIT message
- the shared secret  $K$
- the exchange value  $f$
- the exchange value  $e$
- the host public key

and signs it with its private key to generate the signature  $sigH$ .

Those two values, together with the host public key in a raw format, are sent to the client, which verifies the signature and computes the same hash value. Notice that because the public key is sent in a raw format, the client must have a copy of the server's public key stored locally, and this is one of the major attack vectors of SSH.

Notice also that, because many algorithms need an initialization vector, which is computed using an hash algorithm over  $H$ , which makes it a keyed digest, again.



## Key derivation

Many keyed digest require a initialization vector.

Starting with those exchanged values, if there is any kind of algorithm that requires an IV (nearly all), then the initial IV is:

- From client to server:  $HASH(K \parallel H \parallel "A" \parallel \text{session\_id})$ , where  $K$  is the key that has been negotiated and "A" is simply a letter.
- From server to client:  $HASH(K \parallel H \parallel "B" \parallel \text{session\_id})$ .

The encryption key is generated as follows:

- From client to server:  $HASH(K \parallel H \parallel "C" \parallel \text{session\_id})$ .
- From server to client:  $HASH(K \parallel H \parallel "D" \parallel \text{session\_id})$ .

The integrity key:

- From client to server:  $HASH(K \parallel H \parallel "E" \parallel \text{session\_id})$ .
- From server to client:  $HASH(K \parallel H \parallel "F" \parallel \text{session\_id})$ .

The HASH algorithm is the one used to create all these keys, and this could be a weak point because it would be much better to generate keys using a KDF (Key Derivation Function).

### 2.1.2 Encryption

Encryption is compulsory in SSH, which algorithms are negotiated during the key exchange. The supported algorithms are:

- (required for backward compatibility) 3des-cbc (w/ three keys, i.e. 168 bit key)
- (recommended) aes128-cbc
- (optional) blowfish-cbc, twofish256-cbc, twofish192-cbc, twofish128-cbc, aes256-cbc, aes192-cbc, serpent256-cbc, serpent192-cbc, serpent128-cbc, arcfour, idea-cbc, cast128-cbc, none

The key and initialization vector needed are established during the key exchange and all packets sent in one direction is a single data stream, meaning that the first packet sent has got an initialization vector, which is the one that was created with the key derivation function. But in the following ones, the IV is passed in the packet from the end of one to the beginning of the next one. In fact, there are two cases:

- if **CBC mode** is used, the IV for next packet is the last encrypted block of the previous packet
- if **CTR mode** is used, the IV for the next packet is the incremented value of the last IV

### 2.1.3 MAC

The MAC algorithm and the key are negotiated during the key exchange and they can be different in each direction.

The basic set of algorithms are:

- hmac-sha1 (required for backward compatibility) [key length = 160-bit]
- hmac-sha1-96 (recomm) [key length = 160-bit]
- hmac-md5 (opt) [key length = 128-bit]
- hmac-md5-96 (opt) [key length = 128-bit]

The mac is computed using the key for the right direction over the sequence number concatenated with the packet in clear text.

The sequence number is implicit, as in TLS, and its possible because of TCP. It is represented as a 32-bit unsigned integer, which is incremented by one for each packet sent.

This value is never reset, even if keys and algorithms are re-negotiated, which means that when the maximum value is reached, the channel has to be reset.

#### 2.1.4 Peer authentication

The **server** is authenticated using an **asymmetric challenge-response** mechanism, which requires an explicit server signature of the key exchange hash  $H$ . (The challenge is implicit, the signature is implicit).

The client locally stores the public keys of the server, which are usually stored in the `known_hosts` file in the `.ssh` directory.

When connecting to a server not listed in that file, the client is asked to store the public key of the server in the file, following a **TOFU** (Trust On First Use) policy.

Because of this, a good practice is to protect the `known_hosts` file for authentication and integrity, and to periodically audit/review all the `known_hosts` files to quickly detect added/deleted hosts or changed keys.

Client authentication can be performed by two methods. The first one is based on **credentials** (username and password), which are exchanged only after the protected channel is established. This method protects against sniffing attacks, but not other ones like password guessing.

The second method is based on **asymmetric challenge-response**. In that case, the server must locally store the public keys of the users allowed to connect, typically in the `authorized_keys` file in the `.ssh` directory.

As per the server authentication process, as a good practice, the `authorized_keys` file should be protected for authentication and integrity, and periodically audited/reviewed to quickly detect added/deleted keys or changed keys.

## 2.2 Port forwarding

We already mentioned that SSH can act as a site-to-site VPN, but rather than creating a general purpose VPN, SSH can be used to perform **port forwarding** or **tunneling**.

Port forwarding is a technique that allows to forward unprotected TCP traffic through a secure SSH channel.

This allows to secure unprotected service like POP3, SMTP or HTTP while also allowing the application to perform their normal authentication over an encrypted channel.

There are two types of port forwarding, **local** and **remote**, but both use the Connection Protocol to encapsulate a TCP channel inside a SSH one.

### 2.2.1 Local port forwarding

The concept of local port forwarding is to forward a local port to a remote one, which means that the client listens on a local port and forwards the traffic to a remote port.

For example, let's suppose to be behind a firewall that blocks access to an external mail server because only secure traffic is permitted. An SSH tunnel can be created to forward the local port 110 to the remote port 25 using the following command:


```
ssh -L 1234:mail_server:25 user@ssh_server
```

the traffic to port 1234 on the (internal) client will be forwarded to port 25 on the mail\_server



by using a tunnel to the (external) `ssh_server`.

The firewall will see the traffic as SSH traffic, but the user has still to configure the mail user agent to use the local port 1234 to connect as the outgoing mail server.



### 2.2.2 Remote port forwarding


Remote port forwarding is the opposite of local port forwarding, which means that the client listens on a remote port and forwards the traffic to a local port.

For example, let's suppose to have a web server on a local machine that is behind a firewall, and the user wants to make it available to the public internet. The user can create a tunnel to the remote machine with the following command:

```
ssh -R 8000:localhost:80 user@ssh_server
```

The traffic to port 8000 on the remote machine will be forwarded to port 80 on the local machine.

The user can now access the web server on the local machine by connecting to the remote machine on port 8000.



### 2.2.3 Causes of insecurity

In general, the biggest causes of insecurity for SSH are direct trust in the public keys. In fact, x.509 certificates are not used, but some commercial SSH implementations support them and openSSH has SSH certificates, which are a different thing.

Furthermore, blindly trusting the public keys can lead to a MITM attack. The server could also have a weak platform security (worms, malicious software, rootkits, etc.) as well as the client (malware, keyloggers, etc.).

One last point is that any connection to a local forwarded port will be tunneled, even if coming from another node. This is undesirable, as such it is recommended to specify the binding address when creating a tunnel:

```
ssh -L 127.0.0.1:1234:mail_server:25 user@ssh_server
```

## 2.3 Attacks on SSH

### 2.3.1 BothanSpy

BothanSpy is believed to be a CIA tool, according to information released by Wikileaks on July 6, 2017, as part of their Vault 7 disclosure. This malware targets Xshell, a Windows-based SSH client widely used in the USA and South Korea. BothanSpy operates by injecting a malicious DLL into the Xshell process, enabling it to steal sensitive information from SSH connections. Specifically, it collects:

- Usernames and passwords from password-authenticated connections.
- Usernames, private key filenames, and passphrases from public-key authenticated connections.

BothanSpy is designed to work with the ShellTerm attack framework, featuring:

- A covert communication channel with a Command and Control (C&C) server.
- DLL injection capabilities for direct C&C communication, with no data written to disk, making it difficult to detect using traditional anti-malware solutions.
- An offline mode, where the collected data is written to disk, encrypted with AES for later retrieval.

### 2.3.2 Gyr Falcon

Gyr Falcon is another suspected CIA tool, as revealed by Wikileaks on July 6, 2017, in their Vault 7 release. This malware specifically targets OpenSSH on enterprise Linux systems, such as RedHat and CentOS. Gyr Falcon operates by pre-loading a malicious DLL to intercept plaintext traffic, capturing information both before encryption and after decryption. The tool is capable of capturing:

- Usernames and passwords.
- Actual data transmitted over SSH connections.

Key features of Gyr Falcon include:

- An encrypted configuration file.
- An encrypted file for storing captured data.

While Gyr Falcon requires root access for installation, it notably does not integrate with a Command and Control (C&C) server. Instead, it freely reads and writes files on disk, potentially exploiting the limited use of anti-malware solutions on Linux platforms. Despite these capabilities, Gyr Falcon is not particularly stealthy or sophisticated compared to other tools.

### 2.3.3 Brute force attack

Brute force attacks can exploit the false sense of security provided by secure communication channels, especially when insecure authentication methods like reusable passwords are employed. In a typical brute-force attack, an attacker systematically attempts to guess a password by trying all possibilities from a dictionary of commonly used passwords.

An illustrative example of such an attack occurred in September 2016:

- Two servers, one with an IPv4 address and the other with an IPv6 address, were activated in the cloud to test SSH brute-force attacks.
- The IPv4 server was immediately targeted and compromised within just 12 minutes, primarily because the root password was set to "password."
- Once compromised, the server was quickly used in a Distributed Denial of Service (DDoS) attack.
- In contrast, after one week, the IPv6 server remained unscathed, as it had not even been attacked.

This example demonstrates the risks of relying on weak passwords and highlights the disparity in attack frequency between IPv4 and IPv6 networks.

#### Protection from brute-force attacks

To mitigate the risk of SSH brute-force attacks on Linux systems, several strategies can be implemented:

- **Account Lockout:** Automatically lock accounts after a certain number of failed authentication attempts using tools like `pam_tally2` or `pam_faillock`.
- **TCP Wrappers:** Use TCP wrappers to restrict SSH access by permitting or denying connections from specific hosts or networks.
- **SSH Rate Control:** Implement rate limiting for SSH connections using IPtables, for example, by allowing only 5 connection attempts per minute.
- **Disable Root Access:** Deny direct SSH access to the root account to limit exposure.
- **Change Default SSH Port:** Move SSH from the default port 22 to another port to reduce the chances of automated scans and attacks.
- **Fail2ban:** Use `Fail2ban` to monitor log files and blacklist IP addresses after a defined number of authentication failures.
- **Two-Factor Authentication (2FA):** Implement 2FA, such as Google Authenticator, for an added layer of security.
- **Disable Password-Based Authentication:** Disable password-based authentication altogether in favor of more secure methods like public-key authentication.

These measures can significantly reduce the effectiveness of brute-force attacks and improve the overall security of SSH services.

## 2.4 Main Applications of SSH

SSH (Secure Shell) is widely used for a variety of secure applications due to the encryption and authentication mechanisms it offers. The main applications include:

- **Remote Interactive Access:** SSH allows users to securely connect to remote systems with a text-based interface, enabling command-line interaction.
- **Remote Command Execution:** SSH facilitates the secure execution of commands on remote systems, making it ideal for system administration and automation tasks.
- **Tunneling:** SSH can create secure tunnels for various applications, such as forwarding network traffic, encrypting data, or creating virtual private networks (VPNs).

Keep in mind that SSH is now natively available on Linux, macOS, and Windows, providing both client and server functionality on these platforms.

All of these functions benefit from the security properties inherent to SSH, including confidentiality, integrity, and authentication.

## Chapter 3

# The X.509 standard and PKI

We will now discuss the X.509 standard and Public Key Infrastructure (PKI), which is widely adopted for public key certification.

### 3.1 Public-key certificate (PKC)

Lets begin with a definition:

A **public-key certificate** is a **data structure** to securely bind a public key to some attributes.

A PKC is said to "securely bind" a public key to some attributes because it is signed by a trusted entity, the **Certification Authority**, but other techniques for certification exists (e.g. blockchain, direct trust and personal signature).

The attributes in the certificate are those employed in the transaction being protected by the PKC, which are difficult to decide a-priori without knowing the context: for example one attribute could be an email address associated with the entity, which may not be valid for all the transactions.

PKC are most important to achieve **non-repudiation**, many certificates have legal value, and **digital signature** in a secure way.

Keep in mind that a PKC is the complement of a corresponding **personal private key**, which is kept secret by the entity.

Which means that if the private key is compromised, the PKC is compromised as well.

#### 3.1.1 Key Generation in Public Key Cryptography (PKC)

Key generation in public key cryptography (PKC) involves complex algorithms and often relies on random number generators (RNGs) to ensure strong, unpredictable keys. Once a key is generated, the private key must be securely protected in various contexts:

- **Storage:** The private key needs to be securely stored to prevent unauthorized access.

- **Usage:** The private key must also be protected when being used, as it could be exposed during operations in the CPU as it is used in clear text.
- **Software Application:** In software environments (e.g., web browsers), the trustworthiness of the computing platform must be considered, as it may be vulnerable to malware or weak implementations. This is true both for the context of key generation and use.
- **Dedicated Hardware:** Storing and using keys in dedicated hardware, such as a smart card, offers enhanced protection but comes with limitations in terms of algorithm updates or vulnerability patching( which may be difficult or even impossible). For example many old smart cards use 1024-bit RSA keys, which are considered insecure today.
- **Key Injection:** Keys may be generated in software and injected into hardware devices, a process that can be useful for key recovery, but should be restricted to encryption keys to maintain security and ensure non repudiation. This is most important for the private key.

PKC key management requires careful consideration of both security and operational challenges, especially when dealing with long-term security mechanisms.

### 3.1.2 Certification architecture

In Public Key Infrastructure (PKI), several key entities are responsible for managing the lifecycle of Public Key Certificates (PKCs):

- **Certification Authority (CA):** The CA is responsible for generating and revoking PKCs. It also publishes PKCs and maintains information about their status, such as whether they are valid or revoked, or even suspended.
- **Registration Authority (RA):** The RA plays a critical role in verifying the claimed identity and attributes of a certificate requestor. It authorizes the issuance or revocation of PKCs but does not generate them itself.
- **Validation Authority (VA):** The VA provides services that allow third parties to verify the validity status of a PKC, because some responsibilities or the RA may be delegated to it. This may involve downloading Certificate Revocation Lists (CRLs) or querying an Online Certificate Status Protocol (OCSP) responder to check the certificate's status in real-time.
- **Revocation Authority:** Although not an official term, this role may be assigned to either the RA or CA. The revocation authority is delegated to revoke certificates, often on a more urgent basis than their issuance (they are available most of the time), ensuring that compromised or invalid certificates are promptly rendered unusable.

These entities work together to ensure the security and trustworthiness of PKC-based systems by handling key tasks such as certificate issuance, verification, and revocation.

### 3.1.3 Certificate generation

The general When an actor want to generate a certificate, it must first generate a key pair (public and private key). The private key is stored locally and protected(encrypted) as good as possible, while the public key is sent to the CA with some attributes that help identify the actor.

The CA requires that the attributes are both correct and are associated actually with the entity that is in control of the private key, which requires authentication of the actor. For that purpose, if the entity is a human it is usually required to go physically to the CA, even if in some cases a video call may be enough, with the requirement of showing a valid ID.

The Registration Authority will verify the attributes and the identity of the actor, and if everything is correct, it will send a message to the CA to generate the certificate.

At this point the CA generates the certificate, signs it with its private key and sends it back to the actor while also storing it in its repository.

As you may noticed, the verification is only on the attributes, and not of the possession of the private key, which will be discussed later.

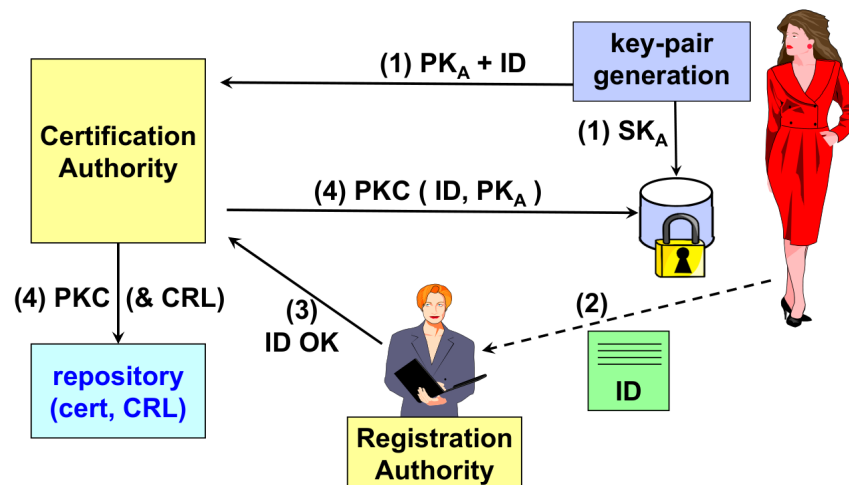


Figure 3.1: Certificate generation

#### Another certificate generation architecture

In addition to what just discussed, alternative architectures for certificate generation exist:

- **Key Pair Generation by RA:** In some architectures, the Registration Authority (RA) generates the key pair on behalf of the user, obtains the Public Key Certificate (PKC), and securely distributes the key pair and certificate to the user via a secure device, such as a smart card. This approach is often used in large companies where the employees are already known to the organization.
- **User Authentication via Code:** Another common method involves the user visiting the RA to obtain a code that can be used to authenticate their certificate request to the Certification Authority (CA). The code is typically calculated as a Message Authentication Code (MAC) using the shared secret key between the RA and CA:

$$\text{code} = \text{MAC}(K, ID)$$

where  $K$  is a shared secret key between the RA and CA, and  $ID$  is the user's identity. The user then submits this code to the CA to validate the certificate request.

These alternative methods provide flexibility for organizations with different security needs, allowing for centralized key generation and secure certificate distribution.

## 3.2 X.509 certificates

X.509 is an ITU-T (international telecommunication union body) standard that defines the format of public key certificates used to verify the identity of a key owner in cryptographic systems. X.509 was developed as a certification technology to protect the x500 directory service, and has undergone several versions:

- **v1 (1988)**: The initial version of the standard.
- **v2 (1993)**: A minor update with small improvements.
- **v3 (1996)**: Added extensions and introduced version 1 of the attribute certificate.
- **v3 (2001)**: Further enhancements, including version 2 of the attribute certificate, which are not used anymore.

X.509 is part of the larger X.500 standard for directory services, often referred to as "white pages," which are used for managing information about entities in a structured way (directory services). X.509 provides a solution to the problem of securely identifying the owner of a cryptographic key.

The certificates and their structures are defined using **ASN.1** (Abstract Syntax Notation One), a standard interface for defining data structures exchanged in networking environments in a neutral and platform-independent way.

### 3.2.1 PKC Scope

A Public Key Certificate (PKC) contains information that uniquely associates a cryptographic key with an entity. This binding between the key and the entity is ensured by a **Trusted Third Party (TTP)**, typically referred to as a Certification Authority (CA), which digitally signs each certificate to guarantee its authenticity.

The liability associated with a PKC may be restricted to specific applications or purposes, as outlined in the CA's certification policy. This policy defines the intended usage and limits the scope of the certificate, ensuring it is applied within the proper legal and technical contexts.



### 3.2.2 Certificate Policy (CP) and Certification Practice Statement (CPS)

According to RFC-3647, "Internet X.509 Public Key Infrastructure Certificate Policy and Certification Practices Framework," both the Certificate Policy (CP) and Certification Practice Statement (CPS) play key roles in defining the use and management of Public Key Certificates (PKCs):

- **Certificate Policy (CP):** A CP is a named set of rules that defines the applicability of a PKC to a specific community or class of applications with common security requirements. It establishes the minimum requirements for the issuance and management of certificates and can be followed by multiple Certification Authorities (CAs). For example, a government CP may apply to all certification providers issuing certificates for official use.
- **Certification Practice Statement (CPS):** A CPS outlines the specific practices that a CA follows when issuing PKCs. While a CP specifies the general rules, the CPS provides detailed implementation procedures. Each CA develops its own CPS, which is tailored to its operations and describes how it meets the requirements set forth in the CP.

### 3.2.3 X.500 Directory Service

The X.500 directory service was the first system to use X.509v1 certificates. It aimed to manage information about people and entities in a network. However, this early setup had three big issues:

- **No Guarantees on the CA's Quality:** There weren't clear rules or policies to ensure that Certification Authorities (CAs) were reliable or trustworthy. People just had to trust the CA, which caused concern over the security of the certificates.
- **Lack of a proper X.500 Infrastructure:** Even though certificates were supposed to be stored in X.500 directories, the infrastructure to do this was never fully implemented worldwide. This made it tough to access certificates when needed.
- **Hard to Establish Certification Paths:** It was difficult to connect two users with certificates issued by different CAs because the relationships between CAs weren't well-defined. With each CA running its own domain, figuring out how to establish a secure connection between users from different domains was tricky.

These issues slowed down the adoption of X.500 and made it clear that better systems were needed for managing certificates.

#### Fixing X.509v1 Problems

To solve these problems, X.509v1 was updated with a couple of key changes:

- **Move the semantics Outside the Certificate:** Instead of relying on the certificate itself to define its semantics, they decided to put that responsibility on the application

or some external system. This approach was used in things like RFC-1422 (PEM), where the application takes care of interpreting the certificate.

- **Make Certificates More Flexible:** The original X.509v1 certificates were pretty limited: they only had an identifier and a key. With X.509v3, they added more fields and options to make certificates more flexible and useful for a variety of purposes. This allowed them to handle a wider range of security scenarios.

These updates helped fix the early issues and set the stage for broader adoption of the improved X.509v3 certificates.

### 3.2.4 RFC-1422

RFC-1422 tried to fix some of the problems with the early X.500 systems by proposing a worldwide certification hierarchy. The plan was to have one main CA at the top: the **IPRA** (Internet Policy Registration Authority). This CA would be responsible for setting the rules and policies for all other CAs in the system, ensuring that everyone followed the same security practices.

Instead of directly issuing certificates to lower-level CAs, the IPRA would issue certificates to **Policy Certification Authorities (PCAs)**.

The IPRA oversaw the entire structure and laid out specific roles for different types of Certification Authorities (CAs) to manage and enforce certificate policies:

- **Policy Certification Authority (PCA):** PCAs were responsible for setting and enforcing the policies that governed how certificates were issued. They made sure that the CAs under them followed the right security practices.
- **Name Subordination:** CAs had to issue certificates within a defined part of the naming structure, ensuring a consistent and organized system. For example:
  - **CA n.1:** C=IT (country-level CA for Italy)
  - **CA n.2:** C=IT, O=Politecnico di Torino (CA for an organization within Italy)

However, this system wasn't without flaws—it was still based on geographic hierarchies, which didn't always match the needs of modern, global organizations.

The IPRA was established, and four PCAs were created under it. One of these PCAs was designated as **high assurance**, meaning it had to carry out stricter checks before issuing certificates. For instance, BBN, a company heavily involved with the U.S. military, might require things like fingerprints or even DNA tests to verify identities.

In contrast, **mid-level assurance** certificates, like those issued by universities such as MIT, required less rigorous checks—maybe just showing a passport. MIT could even set up its own subordinate CAs, like one for the Laboratory for Computer Science, to handle its internal certificate needs.

What's interesting is how different cultural practices shaped the certification process. In the U.S., where people don't generally carry ID cards, **residential certification authorities** were used. If you wanted to open a bank account, for example, you might have to show an electricity bill to prove your address. While this seems strange in countries with formal ID systems, in the U.S., it was a necessity.

Finally, there were **persona certificates**, which allowed for anonymous certificates. These were for users who didn't want to reveal their identity but still needed secure communications. For example, you could be "anonymous user number 95." Even though no one knew who you were, the security of your communication was still guaranteed.




Figure 3.2: Internet PEM hierarchy (RFC-1422)

Unfortunately, this system was a complete failure for several reasons:

- The hierarchical structure severely limited flexibility, much like the issues we saw with X.500. For international companies, this rigid hierarchy simply didn't work.
- Name subordination imposed additional restrictions. You couldn't just assign any name—you had to follow the strict hierarchy, which wasn't always practical.
- The use of the PCA concept lacked flexibility, especially in commercial settings. Instead of automating the process, a human operator was needed to check if the requester was following the policy. This made the system cumbersome and inefficient for businesses.
- Lastly, the biggest issue was trust. Where would the IPRA be placed? Naturally, the U.S. wanted it, but other countries, like Russia, China, or even Japan, wanted control as well. No one trusted a single country to sit at the top of the hierarchy, fearing the possibility of a fake hierarchy being created. As a result, the experiment collapsed. It briefly existed, primarily used by the United States, but it never gained global traction.

### 3.3 X.509 Version 3

X.509 Version 3 is a standard that was finalized in June 1996 through a collaborative effort between ISO, ITU, and the IETF, with the goal of making certificates suitable for internet applications. This version consolidated all the necessary updates to certificates and Certificate Revocation Lists (CRLs) into a single document. Earlier versions of X.509 (v1 and v2) relied on external semantics, such as the failed attempt with the Internet PEM hierarchy, which proved unsuccessful. X.509v3 addressed these shortcomings by ensuring certificates would be useful for internet applications, particularly as OSI applications became

mostly obsolete.

Key features of X.509 Version 3 include:

- **Types of Extensions:**

- **Public Extensions:** These are defined by the standard and made publicly available for anyone to use. Because they are standardized, any system or application should be able to understand and process them.
- **Private Extensions:** These are tailored to specific user communities and are not publicly available. If a system does not understand a private extension, it will treat it as a binary blob and discard it. Private extensions are crucial for certain organizations that require custom functionality.

The introduction of extensions is a major improvement in X.509v3, adding flexibility without changing the fundamental structure of the certificate (which remains the same as in X.509v1). Extensions are contained within a small field but open up a wide range of possibilities. Public extensions ensure standard compatibility, while private extensions allow customization specific to organizations.

- **Certificate Profile:** This refers to a set of extensions tailored for a specific purpose or application, ensuring that certificates meet particular requirements. For example, RFC 5280, titled "Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile," provides guidelines on how X.509 certificates and CRLs should be structured for protecting internet applications. Certificate profiles are important because they dictate how certificates should be used for different scenarios, ensuring consistency and security.

### 3.3.1 Base syntax

The base syntax of an X.509 certificate is defined as follows:

```
Certificate ::= SEQUENCE {
    signatureAlgorithm AlgorithmIdentifier,
    tbsCertificate TBSCertificate,
    signatureValue BIT STRING
}

TBSCertificate ::= SEQUENCE {
    version [0] Version DEFAULT v1,
    serialNumber CertificateSerialNumber,
    signature AlgorithmIdentifier,
    issuer Name,
    validity Validity,
    subject Name,
    subjectPublicKeyInfo SubjectPublicKeyInfo,
    issuerUniqueID [1] IMPLICIT UniqueIdentifier OPTIONAL,
    -- if present, version must be v2 or v3
}
```

```

    subjectUniqueID [2] IMPLICIT UniqueIdentifier OPTIONAL,
        -- if present, version must be v2 or v3
    extensions [3] Extensions OPTIONAL
        -- if present, version must be v3
}

```

This abstract notation outlines the basic structure of the certificate. The **Certificate** is defined as a sequence, which is a data structure that includes several fields. The first field, **signatureAlgorithm**, is the algorithm identifier used by the Certificate Authority (CA) to sign the certificate.

Next, the **tbsCertificate** (to be signed certificate) is of type **TBSCertificate**, which contains the main information that needs to be signed. Finally, the **signatureValue** is a bit string that holds the actual signature.

The **TBSCertificate** is again defined as a sequence, starting with the **version** field, which begins numbering with zero, indicating version 1. For version 3, this field has a value of two. Following that, we have the **serialNumber** of the certificate and the **signature** algorithm identifier.

You may wonder why the signature algorithm appears twice—once externally and once within the **TBSCertificate**. The external algorithm identifier is crucial for verifying the signature; without knowing the algorithm, you cannot validate the signature value. Having it specified internally helps ensure that the signature algorithm used for verification matches the algorithm used for signing. If they differ, it's a red flag, indicating that the certificate should not be trusted.

The **issuer** field identifies the certification authority by name, and the **validity** field specifies the validity period of the certificate. The **subject** is identified by its name, and the **subjectPublicKeyInfo** field contains both the algorithm and the actual public key of the subject.

In addition, versions 2 and 3 introduce optional unique identifiers for both the issuer and subject, though these identifiers are rarely used today. Finally, the **extensions** field, which is optional, must be present for certificates of version 3, as earlier versions do not support extensions.

### 3.3.2 Critical Extensions

In the context of X.509 certificates, extensions can be classified as critical or non-critical, each affecting the certificate verification process differently:

- **Critical Extensions:** If a certificate contains an unrecognized critical extension, it **MUST** be rejected during the verification process. This ensures that any essential information required for proper validation is recognized and handled appropriately.
- **Non-Critical Extensions:** Conversely, a non-critical extension **MAY** be ignored if it is unrecognized. This allows flexibility in certificate processing, as non-critical information does not impede the overall verification of the certificate.

The responsibility for handling these extensions lies entirely with the entity performing the verification, referred to as the **Relying Party (RP)**. The RP must implement logic

to correctly interpret and respond to both critical and non-critical extensions in accordance with their definitions.

### 3.3.3 Public Extensions

X.509 version 3 defines four classes of public extensions that enhance the functionality and applicability of certificates:

- **Key and Policy Information:** Extensions in this class provide additional information regarding the key usage and policies applicable to the certificate, guiding how the key should be used within specific contexts.
- **Certificate Subject and Certificate Issuer Attributes:** These extensions include attributes related to both the subject of the certificate (the entity that the certificate represents) and the issuer (the Certification Authority), enabling more detailed identification and classification.
- **Certificate Path Constraints:** These extensions define rules and limitations regarding the certification path, ensuring that certificates can only be used in certain contexts or under specific conditions, enhancing security within the certificate hierarchy.
- **CRL Distribution Points:** This extension indicates where the Certificate Revocation List (CRL) can be found, providing necessary information for relying parties to check the revocation status of the certificate efficiently.

### 3.3.4 Key and policy information

There are six public extensions dedicated to this purpose. We won't discuss all of them, but we will focus on the most important ones.

#### Authority Key Identifier

Key and policy information extensions in X.509 v3 provide crucial details about the public keys associated with certificates. One significant component is the **Authority Key Identifier (AKI)**, which serves the following purposes:

- **Identification of the Signing Key:** The AKI identifies a specific public key used to sign a certificate, ensuring that the verification process can accurately trace the certificate's authenticity back to the correct authority.
- **Identification Methods:** The identification can be achieved through:
  - A **key identifier**, typically represented as the digest (hash) of the public key of the certification authority.
  - The combination of **issuer-name** and **serial-number**, allowing a clear reference to the issuing CA's key.
- **Usage:** The AKI is particularly useful in scenarios where the same CA might utilize multiple keys, such as for different assurance levels, like low assurance for general use and high assurance for sensitive applications.

- **Non-Critical Extension:** While the AKI is classified as non-critical, its presence can be vital in certain applications. In the professor experience, the absence of this extension can lead to many applications rejecting the certificate as invalid.
- **Importance in Certificate Chains:** When an application receives a certificate, it often needs to establish a trust path up to a root CA. This path-building process relies heavily on the AKI field. If a CA neglects to include this extension, it may lead to compatibility issues with various applications, despite being non-critical.

### Subject key identifier

The **Subject Key Identifier** (SKI) extension serves to identify a specific public key associated with a certificate. Key features include:

- **Identification of Public Key:** The SKI uniquely identifies a particular public key used in an application. This is especially important in scenarios where the public key may be updated or replaced over time.
- **Being Non-Critical:** The SKI is classified as a non-critical extension, meaning that while it provides valuable identification information, its absence does not necessarily prevent the certificate from being considered valid.

This one is usually not present in the certificate because an identifier is already present in the certificate.

### Key usage

The **Key Usage** (KU) extension specifies the application domains in which a public key may be utilized. Key characteristics include:

- **Application Domain Identification:** The KU extension identifies the specific purposes for which the associated public key can be employed, ensuring that it is used appropriately within defined contexts. In case of misuse, the application may refuse to accept the certificate and the CA may refuse liability.
- **Critical or Non-Critical:** The Key Usage extension can be classified as either critical or non-critical:
  - If marked as **critical**, the certificate may only be used for the specific purposes indicated in the Key Usage extension. Any usage outside the defined scopes would render the certificate invalid.
  - If marked as **non-critical**, the certificate can be used more flexibly, potentially allowing usage beyond the defined applications without invalidating the certificate.
- **Permitted Cryptographic Operations:** The following cryptographic operations can be defined within the Key Usage extension:
  - **digitalSignature:** valid in the certificate of a user but also in the certificate of a certification authority.

- **nonRepudiation**: Specifically permitted for users (non repudiation is only for humans after all).
- **keyEncipherment**: Permitted for users.
- **dataEncipherment**: Allowing encryption of data.
- **keyAgreement**: Involves:
  - \* **encipherOnly**: Restricting use to enciphering operations.
  - \* **decipherOnly**: Restricting use to deciphering operations.
- **keyCertSign**: Specifically permitted for Certificate Authorities (CAs).
- **cRLSign**: Also permitted for Certificate Authorities (CAs) for signing Certificate Revocation Lists (CRLs).

### Private key usage period

The **Private Key Usage Period** extension in X.509 v3 specifies the time frame during which the associated private key may be used. Key features include:

- **Usage Period Definition**: This extension defines the period during which the private key can be actively utilized, helping to enforce time-based restrictions on key usage.
- **Non-Critical Extension**: The Private Key Usage Period extension is always classified as non-critical, meaning that its absence does not invalidate the certificate. However, the information it provides can be important for managing key lifecycles.
- **Usage Discouraged**: While this extension is available, its use is generally discouraged in practice. Organizations often prefer to manage key lifetimes through other means, such as regular key rotation, rather than relying on a specified usage period within the certificate.

### Certificate policies

The **Certificate Policies** extension outlines the specific policies that were adhered to during the issuance of the certificate and defines the purposes for which the certificate can be utilized. Key aspects include:

- **Policy Listing**: This extension provides a comprehensive list of the policies that govern the certificate issuance process, ensuring clarity regarding its intended use.
- **Indication Methods**: Certificate policies can be indicated using various formats, including:
  - **Object Identifier (OID)**: A unique identifier for the policy.
  - **Uniform Resource Identifier (URI)**: A link to a location where the policy can be reviewed.
  - **Text Message**: A textual description of the policy.
- **Critical or Non-Critical**: The Certificate Policies extension can be classified as either critical or non-critical:



- If marked as **critical**, the certificate must only be used in accordance with the specified policies; otherwise, it may be considered invalid.
- If marked as **non-critical**, the certificate can be used more flexibly, although the policies still provide guidance for its intended use.
- **Support for Authentication and Authorization:** The use of this extension can enhance not only the authentication of users and entities but also facilitate authorization processes, providing a clearer understanding of the permissible actions associated with the certificate.

## Policy mappings

The **Policy Mappings** extension in X.509 v3 establishes a correspondence between policies across different certification domains. Key aspects include:

- **Mapping of Policies:** This extension indicates how policies from one certification authority (CA) correspond to policies from another, facilitating interoperability and trust among different certificate frameworks. This can be done automatically, but needs human interaction to be done correctly.
- **Presence in CA Certificates:** The Policy Mappings extension is typically present only in CA certificates, allowing certification authorities to define and communicate relationships between their policies and those of other authorities.
- **Non-Critical Extension:** The Policy Mappings extension is classified as non-critical, meaning its absence does not invalidate the certificate. However, it serves as an important tool for enhancing the understanding and usability of certificates across different domains.

### 3.3.5 Certificate subject and issuer attributes

As previously mentioned, the X.509 v3 standard includes extensions that provide additional information about the subject and issuer of a certificate. Without this extension, the only possible names are x500 ones, which aren't really an identifier.

#### Subject Alternative Name (SAN)

The Subject Alternative Name (SAN) extension provides a flexible way to identify the owner of a certificate. This extension allows the use of various formalisms to represent the owner, including e-mail addresses, IP addresses, and URLs.

Importantly, the SAN field is always considered critical when the subject name field is empty. This means that if there is no subject name, the SAN extension must be present to ensure proper identification and functionality of the certificate.

#### Issuer Alternative Name (IAN)

The Issuer Alternative Name (IAN) extension enables the use of various formalisms to identify the Certificate Authority (CA) that issued a certificate or a Certificate Revocation

List (CRL). This extension supports different identifiers such as e-mail addresses, IP addresses, and URLs.

The IAN field is always considered critical when the issuer name field is empty. In such cases, the IAN extension must be present to ensure proper identification of the issuing authority.

It's also interesting to see what are the possible identifiers that can be used in the IAN field.

One option is the **RFC 802 Name**, which consists of the standard email addresses used in internet applications. The **DNS Name** identifies a node by its domain name, while an **IP Address** identifies the node by its numerical address. Another alternative is the **Uniform Resource Identifier (URI)**, which serves as an entry point on the web. This means that on the same node, you could have different certificates corresponding to various entry points for the applications.

In addition, you can use a **Directory Name**, which can be formatted according to X.500 or LDAP syntax. For example, the LDAP of the Politecnico di Torino might be represented as `country=T; organization=Politecnico di Torino; organizationUnit=Department`. While there is no formal management of country Italy in this case, the syntax remains consistent.

Moreover, **X.400 Addresses** illustrate that these identifiers were still in use when version 3 was developed in 1996, as there were gateways that transformed mail from X.400 to RFC 802 formats. Thus, having the certificate associated with the X.400 address was essential.

Next, we have the **EDI Party Name**, which stands for electronic data interchange. This standard allows companies to avoid manual processing of data. For instance, a car manufacturer might need to check if a supplier has specific tires available. Instead of making a phone call or sending a fax or email, they could use EDI to make a query through a neutral language that facilitates data requests and responses electronically.

EDI serves as a general standard with various implementations, such as EDIFACT, which Stellantis and many other car manufacturers use to manage the availability of necessary items, check prices, and place orders. The significance of EDI is reflected in the certificates associated with it.

Although EDI may seem older compared to modern standards like XML or JSON, many companies still rely on these established systems to avoid dealing with breaking changes that may occur.

In addition to the EDI Party Name, a **Registered ID**, such as a VAT number for a company or an Italian tax number for individuals, can also be used. This represents any official registered identifier for an entity. Another identifier is the **Unique Identifier**, which signifies a distinct identifier for individuals or organizations. Finally, there is the **Other Name** option. This should be used cautiously, as it lacks a well-defined structure. When using the "Other Name," the syntax and content definition fall to the user, making it non-standard and challenging to interpret.

## Subject Directory Attributes

The Subject Directory Attributes extension allows for the storage of specific directory attributes related to the owner of the certificate. For example, the Department of Defense (DoD) utilizes this extension to indicate attributes like "citizenship."

The actual usage of Subject Directory Attributes heavily depends on the application since there are no standard definitions for these attributes. This extension is classified as non-critical, meaning its absence does not prevent the certificate from being valid but may limit its utility in certain contexts.

### 3.3.6 Certificate path constraints

Let's now go over the chain of trust and the constraints that can be created in the certificate path.

There are three main constraints that can be set in the certificate path.

#### Basic Constraints

The correct term for user in the context of certificates is **end-entity** (EE).

#### Basic Constraints

The Basic Constraints extension indicates whether the subject of the certificate can act as a Certificate Authority (CA). The values for this extension are defined as follows:

- **BC=true**: This indicates that the subject is a CA.
- **BC=false**: This indicates that the subject is an End Entity (EE).
- Additionally, it is possible to define the maximum depth of the certification tree, but only if **BC=true**.
- This extension can be classified as either critical or non-critical, though it is suggested to always mark this extension as critical.

#### Name Constraints

The Name Constraints extension is applicable only in Certificate Authority (CA) certificates. This extension defines the space of names that can be certified by a CA and follows the same format as the Alternative Names.

Key points regarding Name Constraints include:

At least one specification must be provided between:

- **permittedSubtree**: This serves as a whitelist for names that are allowed.
- **excludedSubtree**: This serves as a blacklist for names that are not allowed.

The whitelist is processed first, ensuring that allowed names take precedence. Caution is advised: An unspecified format in the whitelist (e.g., **directoryName**) is implicitly permitted. This extension can be either critical or non-critical, though it is important to note that there is no non-critical support by Apple.

## Policy Constraints

The Policy Constraints extension is used by a Certificate Authority (CA) to specify constraints that may require an explicit identification by a policy or that inhibit policy mapping for the remainder of the certification path.

This extension can be either critical or non-critical.

### 3.3.7 CRL distribution point

The CRL Distribution Point extension identifies the distribution point of the Certificate Revocation List (CRL) that is to be used in validating a certificate.

The distribution point can be specified in various forms, including:

- **Directory entry**
- **E-mail or URL**

This extension can be either critical or non-critical.

### 3.3.8 Private Extensions

The **Private Extensions** in X.509 v3 enable the creation of extensions tailored to specific user communities, allowing for customized applications within closed groups.

One key aspect of private extensions is their definition. These are custom-defined features intended for use by particular communities or groups of users. This capability provides both flexibility and specificity in the application of certificates, allowing them to meet the unique needs of different organizations or use cases.

Another important consideration is the examples from the Internet Engineering Task Force Public Key Infrastructure (IETF-PKIX), which has defined three notable private extensions for the Internet user community. The first of these is **Subject Information Access**, which provides essential details on how to access data related to the subject of the certificate. The second extension, known as **Authority Information Access**, offers guidance on how to access information pertaining to the certificate authority, ensuring that users can verify and trust the source of the certificate. Finally, the **CA Information Access** extension supplies vital information for accessing resources linked to the certificate authority, facilitating the management and verification of certificates within the relevant community.

#### Subject Information Access

The **Subject Information Access** extension in X.509 v3 specifies a method for obtaining additional information about the owner of a certificate, particularly useful when a directory service is not employed for certificate distribution. This extension allows users to retrieve extra details regarding the certified end entity, which can be important for verifying credentials or understanding the context of the certificate.

The extension typically includes a method, such as HTTP, HTTPS, or LDAP, which outlines the protocol to be used for accessing the required information. Alongside this method, it provides a name indicating the location or address where the information can be

obtained. For instance, it might direct users to a specific web page related to the individual or point to a local directory entry.

It is important to note that this extension is classified as non-critical, meaning it serves primarily as informational. While it enhances the utility of the certificate by providing additional access methods, it does not impact the fundamental validation of the certificate itself.

### Authority Information Access

The **Authority Information Access (AIA)** extension plays a crucial role in X.509 certificates, as it specifies how to access the information and services offered by the certificate authority (CA) that issued the certificate. This extension is applicable to any certificate, whether from a certification authority or an end entity.

AIA contains several important subfields that facilitate various functions. One key subfield is **certStatus**, which may include a URL for an Online Certificate Status Protocol (OCSP) responder. This is particularly useful for checking the real-time status of a certificate, complementing the Certificate Revocation List (CRL) distribution point.

Additionally, AIA may provide pointers for:

- **certRetrieval**: This subfield allows for the retrieval of the certificate itself from a specified location.
- **caPolicy**: It outlines the policies under which the CA operates, offering insight into the trustworthiness of the certificate.
- **caCerts**: This indicates where to find the CA's own certificates, which is essential for building a certification path.

The presence of AIA in a certificate enables users to locate the services provided by the CA. It acts as a guide for retrieving necessary information, such as following a path to an OCSP responder or accessing the CA's certificates. While the AIA extension can be classified as either critical or non-critical, it is typically recommended for scenarios that require real-time status checks or validation. This extension enhances the functionality of certificates by making it easier to access relevant information necessary for establishing trust.

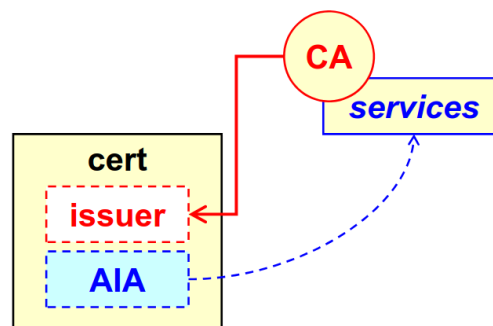


Figure 3.3: Authority Information Access (AIA) extension

## CA Information Access

The **CA Information Access (CAIA)** extension serves a specific function in X.509 certificates, as it indicates how to access information and services provided by the certificate authority (CA) that owns the certificate. Unlike other extensions, CAIA is valid only within a CA certificate itself.

CAIA includes several important subfields that facilitate access to the CA's services, such as:

- **certStatus**: This subfield may point to the status of the CA's certificates, typically through an OCSP responder.
- **certRetrieval**: This allows for the retrieval of the CA's own certificate from a designated location.
- **caPolicy**: This subfield provides information on the policies that govern the CA's operations, helping users understand the trust model.
- **caCerts**: It indicates where the CA's certificates can be accessed, which is vital for validating the certificate chain.

The significance of CAIA lies in its role as a self-pointer for the CA. When you possess a CA certificate and wish to locate its services, CAIA provides the necessary pointers to information about the CA itself, including its policy, OCSP responder, and certificate repository. This self-referential nature distinguishes CAIA from other access methods, which typically point back to the parent CA. As with many extensions, CAIA can be classified as either critical or non-critical, depending on the context in which it is used. However, it is particularly valuable for establishing a complete understanding of the CA's services and trustworthiness.

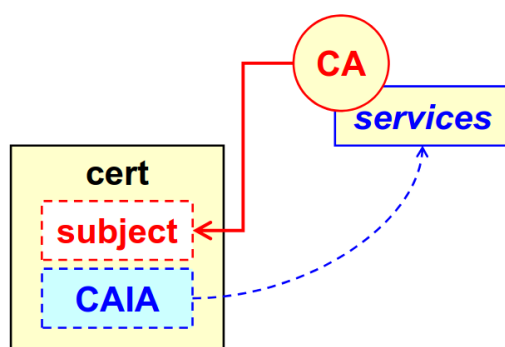


Figure 3.4: CA Information Access (CAIA) extension

## RFC-2459

The **RFC-2459** document defines a profile for the use of X.509v3 certificates in Internet applications, such as IPsec, TLS, and S/MIME. This profile was initially suggested by the Public Key Infrastructure (PKIX) working group to promote interoperability and standardization across diverse systems and protocols.

One of the key aspects of RFC-2459 is that it establishes extensions defined by an Object Identifier (OID), specifically using the base `ID-PKIX ::= 1.3.6.1.5.5.7` (which maps to `iso.org.dod.inet.sec.mechanisms.pkix`). OIDs ensure that each extension and algorithm used within the X.509 framework is uniquely identifiable and properly managed. For instance, PKIX requested an OID rooted in the U.S. Department of Defense namespace due to the historical origin of the internet (as DARPA Net), and further sub-trees were established for protocols like IPsec, TLS, and S/MIME.

RFC-2459 also specifies the supported algorithms that should be implemented to ensure secure and reliable communication. Beyond defining supported algorithms, it offers several implementation suggestions to enhance compatibility. A few examples include:

- **UTC Time:** It is mandatory to include seconds in time fields.
- **Zulu Time:** UTC time must be expressed in the Zulu (Greenwich Mean Time) timezone, which avoids issues with software systems that misinterpret time zones. For instance, some software (notably from Microsoft) might default to local time, causing inconsistencies.
- **Two-Digit Year Format:** Although specifying the year in two digits is not recommended, when used, it should be interpreted within the range of 1950-2049. After 2049, certificates must switch to four-digit year formats to avoid ambiguity and incompatibility issues.

By adhering to the guidelines outlined in RFC-2459, certificate issuers and consumers ensure proper functionality across diverse systems, especially in the context of internet security protocols. The use of standardized extensions, supported algorithms, and implementation details like time format further strengthens interoperability and the reliability of X.509 certificates.

## Extended Key Usage

The **Extended Key Usage** (EKU) extension in X.509 certificates allows for defining specific uses of a certificate in addition to, or in substitution of, the standard `keyUsage` extension. Unlike the standard key usage, which is focused on cryptographic operations, EKU is oriented towards specific applications. This enables a certificate to be more narrowly defined for its intended purpose, enhancing security and clarity in certificate handling.

EKU can be used simultaneously with the standard `keyUsage` extension, though it is important to consider how conflicts between the two might be handled in practice. EKU serves to tailor certificates for specific tasks within certain application domains, which is particularly useful in environments requiring strict certification boundaries, such as server authentication or email protection.

Some of the possible values for EKU include:

- **serverAuth (id-pkix.3.1):** This EKU designates a certificate for server authentication. In cases where serverAuth is used, the equivalent key usages that should be enabled include Digital Signature (DS), Key Encipherment (KE), and Key Agreement (KA).

- **clientAuth (id-pkix.3.2)**: This EKU is intended for client-side authentication and can be mapped to Digital Signature (DS) and Key Agreement (KA).
- **codeSigning (id-pkix.3.3)**: This EKU is used for code signing, allowing developers to sign the code they create. The primary associated key usage is Digital Signature (DS).
- **emailProtection (id-pkix.3.4)**: This EKU provides a certificate for email protection. It includes Digital Signature (DS), Non-Repudiation (NR), Key Encipherment (KE), and Key Agreement (KA) as the relevant key usages.
- **timeStamping (id-pkix.3.8)**: Used for timestamping, this EKU is associated with Digital Signature (DS) and Non-Repudiation (NR). Timestamping certificates are crucial for certifying when a particular action, such as a signature, took place.
- **ocspSigning (id-pkix.3.9)**: This EKU is specific to signing OCSP responses. The associated key usages are Digital Signature (DS) and Non-Repudiation (NR), ensuring that OCSP responses are properly authenticated.

In summary, the PKIX group developed EKU to provide more granular control over certificate usage within specific applications. For example, a certificate intended for email may only be valid for email protection, while another for server authentication may only serve that purpose. The flexibility offered by EKU allows for more precise security policies, tailored to the needs of individual systems or applications.

## Evolution of RFC-2459

The evolution of **RFC-2459**, which initially defined the profile of X.509v3 certificates for use in Internet applications, reflects a series of updates and refinements to both certificate profiles and cryptographic algorithms.

RFC-2459 was replaced by two key documents:

- **RFC-3280**: This document defines the Internet profile of public key infrastructures (PKIs) based on X.509v3 certificates and X.509v2 certificate revocation lists (CRLs). It was intended to standardize how certificates and CRLs are managed and used in Internet-based PKI systems.
- **RFC-3279**: This document covers the algorithms used in conjunction with the RFC-3280 profile, providing identifiers, parameters, and encodings. It includes, and makes obsolete, the earlier RFC-2528, which addressed the use of the Key Exchange Algorithm (KEA).

The decision to split RFC-2459 into two separate RFCs was driven by the recognition that while certificate profiles tend to remain stable over time, cryptographic algorithms evolve more rapidly. By documenting algorithms separately in RFC-3279, it became easier to update the list of supported algorithms without needing to revise the entire certificate profile.

Later on, **RFC-3280** was itself obsoleted by **RFC-5280**, which further refined the Internet PKI profile and introduced updates to align with advancements in cryptographic practices and PKI implementation. RFC-5280 remains a foundational document in the definition of PKI systems used for securing Internet communication.



## RFC-3279 and Supported Algorithms

**RFC-3279** specifies the algorithms that *MUST* be supported by applications using the **RFC-3280** profile. It includes some older algorithms, as well as newer ones. The key algorithms defined in this RFC are:

- **Digest algorithms:**
  - MD-2, MD-5, SHA-1 (preferred)
- **Algorithms for signing certificates/CRLs:**
  - RSA, DSA, ECDSA (Elliptic Curve DSA)
- **Subject public key algorithms (SubjectPublicKeyInfo):**
  - RSA, DSA, KEA(a variant of Diffie-Hellman), DH (Diffie-Hellman), ECDSA, ECDH (Elliptic Curve Diffie-Hellman)

The underlined algorithms were introduced in **RFC-3279** in comparison to **RFC-2459**. One important point is that RFC-3279 permits the use of legacy algorithms, though they are not recommended (e.g., SHA-1 is preferred for digest, and RSA/DSA for certificate signing, with ECDSA added for more modern use).

Once the basic definitions are set, optional algorithms can be defined. Notably:

- **RFC-4055:** Adds better specifications for signatures, including the Probabilistic Signature Scheme (PSS), Optimal Asymmetric Encryption Padding (OAEP), and SHA-2 for digest computation.
- **RFC-4491:** Addresses the needs of Russia by providing specifications for the *GOST* algorithm family, widely used in Russian encryption standards.
- **RFC-5480:** Introduces elliptic curve keys of various kinds.
- **RFC-5758:** Adds support for new algorithms in the *SHA-2* family, such as SHA-224 and SHA-256.
- **RFC-8692:** Uses SHAKE128 and SHAKE256 for signatures, both part of the *SHA-3* family of algorithms.

## RFC-3280 / RFC-5280: PKI Profile

**RFC-3280** and **RFC-5280** specify the profile that defines not only values and fields but also algorithms and procedures. These ensure that certificates are processed consistently across the world. Key aspects include:

- **Path validation algorithm:** Defines how to build or verify the certificate chain, starting from an end entity (EE) certificate, which is issued by a CA, up to a trusted root.
- **Certificate status verification:** Specifies details for verifying the status of a certificate, using methods such as the full CRL or Delta-CRL.

- **Extended Key Usage (EKU):** Adds the *OCSPSigning* extended key usage.
- **Certificate extensions:** Includes important extensions such as:
  - **Inhibit any-policy:** Prevents the use of any policy, ensuring that the CA must specify its own policies.
  - **Freshest CRL:** A pointer to a new type of CRL (Delta-CRL), which contains only the differences with the last full CRL.
  - **Subject information access:** Already discussed previously.
- **Freshest CRL for CRLs:** If there is a Delta-CRL pointer in a certificate or CRL, it is possible to access the differences with the last full CRL. The Freshest CRL is also known as the Delta CRL Distribution Point. It is always marked as non-critical because there are multiple ways to check certificate status (e.g., full CRL, Delta CRL, OCSP), and it is not feasible to make any single method mandatory.

### Additional Points

- **Path validation focus:** Due to the sensitive nature of path validation, recent RFCs have paid much more attention to specifying the algorithms for this process, rather than just providing high-level descriptions as in earlier versions.
- **Support for internationalization:** With global use of certificates, it has become important to support different alphabets, such as those used in Japan, China, and India, in fields like URI, DNS names, and email addresses. **RFC-8398** adds support for internationalized email addresses and domain names.
- **Freshest CRL pointer:** A crucial feature is the *Freshest CRL* pointer, which allows a certificate or CRL to reference the latest Delta-CRL. This prevents the need to redownload a full CRL when only small updates are necessary.
- **No Delta of Delta-CRLs:** Delta CRLs cannot be created incrementally (i.e., you cannot generate a Delta of a Delta CRL). Each Delta CRL must reference the base CRL directly, making the process more straightforward.